

Kinetic Dependence Graphs

M. Amber Hassaan,

Donald Nguyen,

Keshav Pingali

The University of Texas at Austin

Dependence Graph Scheduling

```
A[N,N] = ...
```

```
for i in 1:N
```

```
  for j in 1:N
```

```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```

Application Programmer

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
```

```
  for j in 1:N
```

```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```

Application Programmer



Compiler

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
```

```
  for j in 1:N
```


```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```

Application Programmer



Compiler



Runtime Scheduler

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
```

```
  for j in 1:N
```


```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```

Application Programmer



Compiler



Runtime Scheduler

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
```

```
  for j in 1:N
```


```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```

Application Programmer



Compiler
Parallel Task Library



Runtime Scheduler

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
```

```
  for j in 1:N
```

```
    A[i,j] =
```


```
      A[i-1,j] + A[i,j-1]
```

Application Programmer



Compiler

Parallel Task Library
Parallel Programming Model



Runtime Scheduler

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
```

```
  for j in 1:N
```

```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```

Application Programmer




Compiler

Parallel Task Library



Parallel Programming Model

⋮



Runtime Scheduler


```
A[N,N] = ...  
  
for i in 1:N  
    for j in 1:N  
        A[i,j] =  
            A[i-1,j] + A[i,j-1]
```



Runtime Scheduler

Parallel Task Library
Parallel Programming Model
⋮


9

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
  for j in 1:N
    A[i,j] =
      A[i-1,j] + A[i,j-1]
```

Application Programmer



Compiler

Parallel Task Library
Parallel Programming Model
⋮



Runtime Scheduler


G_0
↑
⋮
 σ_0

Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
  for j in 1:N
    A[i,j] =
      A[i-1,j] + A[i,j-1]
```

Application Programmer




Compiler

Parallel Task Library
Parallel Programming Model
⋮



Runtime Scheduler




Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
  for j in 1:N
    A[i,j] =
      A[i-1,j] + A[i,j-1]
```

Application Programmer




Compiler

Parallel Task Library
Parallel Programming Model
⋮



Runtime Scheduler




Dependence Graph Scheduling

```
A[N,N] = ...
```


```
for i in 1:N
  for j in 1:N
    A[i,j] =
      A[i-1,j] + A[i,j-1]
```

Application Programmer

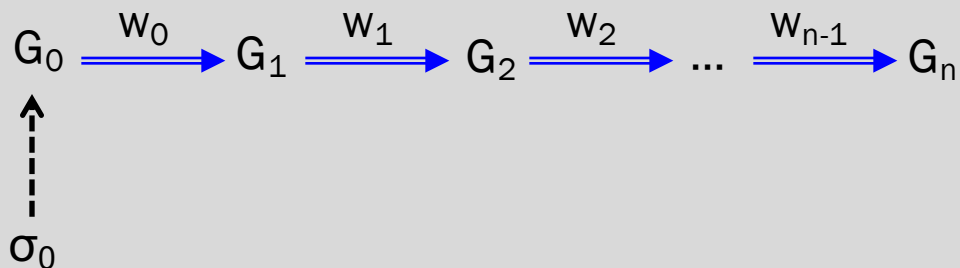


Compiler

Parallel Task Library
Parallel Programming Model
⋮



Runtime Scheduler



Dependence Graph Scheduling



```
A[N,N] = ...
```

```
for i in 1:N
```

```
  for j in 1:N
```

```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```

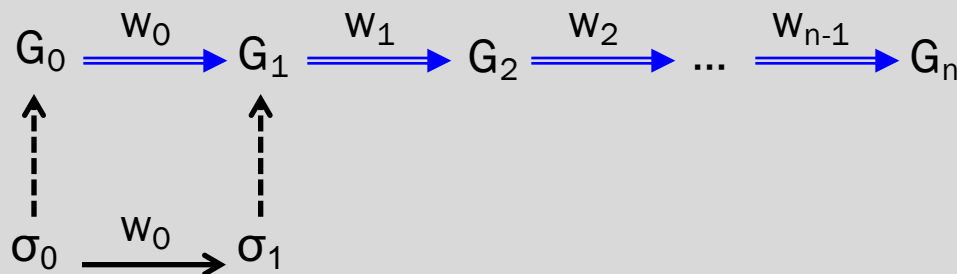


Application Programmer

Compiler

Runtime Scheduler

Parallel Task Library
Parallel Programming Model
⋮



Dependence Graph Scheduling



```
A[N,N] = ...
```

```
for i in 1:N
```

```
  for j in 1:N
```

```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```




Application Programmer

Compiler

Runtime Scheduler

Parallel Task Library
Parallel Programming Model
⋮



Dependence Graph Scheduling



```
A[N,N] = ...
```

```
for i in 1:N
```

```
  for j in 1:N
```

```
    A[i,j] =
```

```
      A[i-1,j] + A[i,j-1]
```



Application Programmer


Compiler

Runtime Scheduler

Parallel Task Library
Parallel Programming Model
⋮


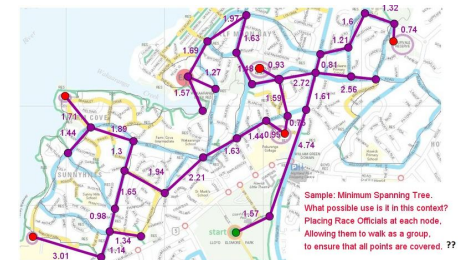

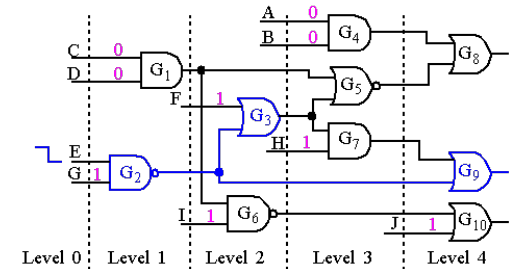
Scheduler
World

Program
World




Limitations of Dependence Graphs


- **Not always sufficient**
 - Discrete event simulation
 - Billiard ball simulation
 - Kruskal's algorithm for MSTs
 - Asynchronous Variational Integrators
 - ...
- **Ordered algorithms**
 - Have tasks with algorithm-specific order in which tasks must appear to execute
- **Can use speculation**
 - But overheads can be high for ordered algorithms




Discrete Event Simulation




Discrete Event Simulation




Discrete Event Simulation




Discrete Event Simulation




Discrete Event Simulation




Discrete Event Simulation




Discrete Event Simulation




Discrete Event Simulation




Discrete Event Simulation





Discrete Event Simulation





Discrete Event Simulation





Discrete Event Simulation




Discrete Event Simulation



Discrete Event Simulation



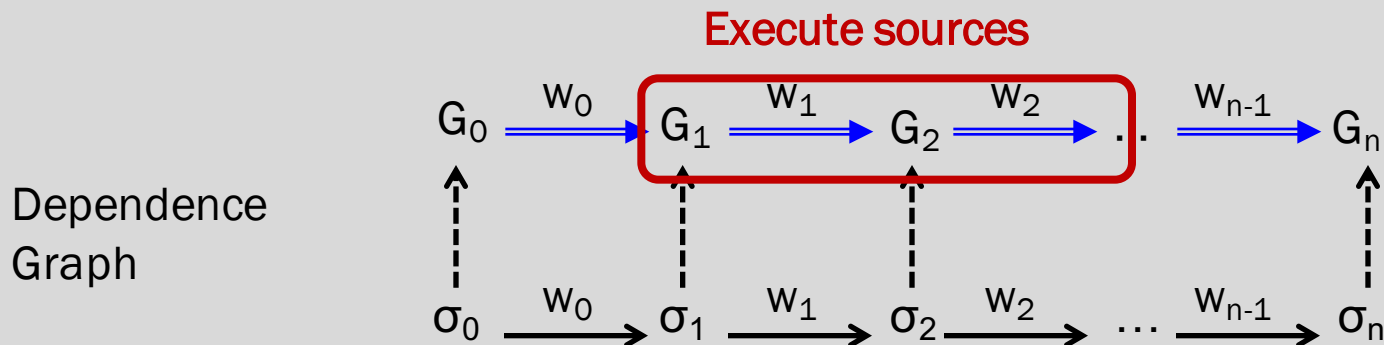
Discrete Event Simulation



- Inadequacy of dependence graphs
 - Tasks are created dynamically
 - Not all sources are safe to execute
 - Task execution may require updating DAG

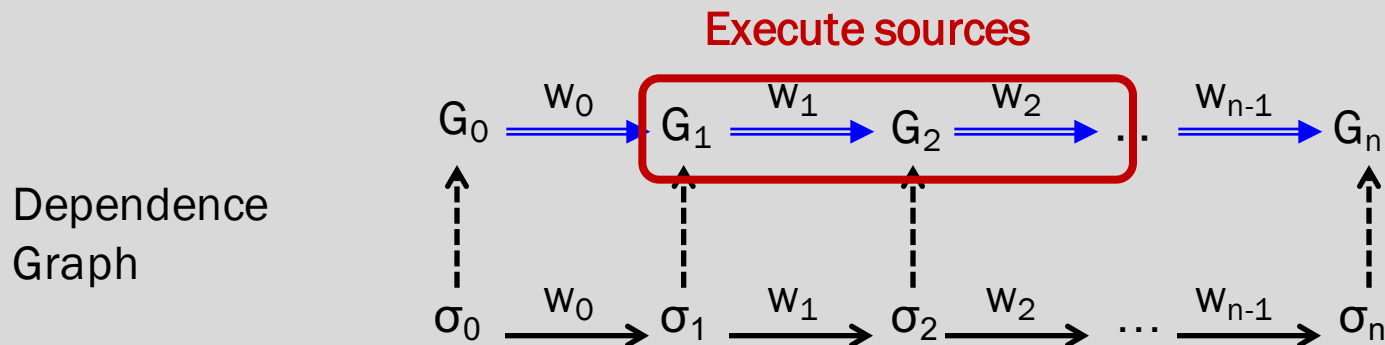
Dependence Graph

- A dependence graph is:
 - A DAG G for the program state σ
 - An **update rule** U to produce the next G after executing task w
 - Remove source



Kinetic Dependence Graph


- A **kinetic** dependence graph is:
 - A DAG G for the program state σ
 - A **safe-source** test P
 - An **update rule** U to produce the next G after executing task w
 - Remove source, update other tasks' dependencies, ...



Kinetic Dependence Graph

- A **kinetic** dependence graph is:
 - A DAG G for the program state σ
 - A **safe-source** test P
 - An **update rule** U to produce the next G after executing task w
 - Remove source, update other tasks' dependencies, ...


Kinetic
Dependence
Graph



Kinetic Dependence Graph


- A **kinetic** dependence graph is:
 - A DAG **G** for the program state σ
 - A **safe-source** test **P**
 - An **update rule** **U** to produce the next **G** after executing task **w**
 - Remove source, update other tasks' dependencies, ...

Kinetic
Dependence
Graph




Safe Sources

- Application-specific cone of influence




Safe Sources

- Application-specific cone of influence




Safe Sources

- Application-specific cone of influence




- In some algorithms, all sources are safe
 - **Stable source** algorithm

General KDG Executor




General KDG Executor


1. Construct initial DAG




General KDG Executor



1. Construct
initial DAG




General KDG Executor




1. Construct
initial DAG



2. Apply safe
source test
3. Execute
4. Update rw-sets
5. Add new tasks



General KDG Executor




a,1


1. Construct
initial DAG




2. Apply safe
source test
3. Execute
4. Update rw-sets
5. Add new tasks




General KDG Executor




1. Construct
initial DAG



2. Apply safe
source test
3. Execute
4. Update rw-sets
5. Add new tasks






General KDG Executor




1. Construct initial DAG



- $a,1$ 
2. Apply safe source test
 3. Execute
 4. Update rw-sets
 5. Add new tasks





General KDG Executor




2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks

1. Construct initial DAG





General KDG Executor




2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks

1. Construct initial DAG





General KDG Executor




2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks



1. Construct initial DAG




General KDG Executor





2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks




General KDG Executor



2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks





General KDG Executor




2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks

1. Construct initial DAG





General KDG Executor




2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks



1. Construct initial DAG






General KDG Executor



2. Apply safe source test
3. Execute
4. Update rw-sets
5. Add new tasks





General KDG Executor




1. Construct
initial DAG




2. Apply safe
source test
3. Execute
4. Update rw-sets
5. Add new tasks






General KDG Executor



1. Construct initial DAG



- $\xRightarrow{\text{a},1}$
2. Apply safe source test
 3. Execute
 4. Update rw-sets
 5. Add new tasks



KDG Specializations

Opportunities

- Each step needs a barrier
- Task's rw-set potentially updated several times prior to execution

KDG Specializations

Opportunities

- Each step needs a barrier
- Task's rw-set potentially updated several times prior to execution

Program Properties

KDG Runtime

Application Programmer

KDG Specializations

Opportunities

- Each step needs a barrier
- Task's rw-set potentially updated several times prior to execution

Program Properties

- Sources
 - **In general**: safe source test could inspect entire state
 - Local safe source test
 - Stable source (sources are safe)
- RW-sets
 - **In general**: task could change other tasks' RW-sets
 - Non-increasing RW-set
 - Structure-based RW-set
- New tasks
 - **In general**: task could create new tasks at any priority
 - Monotonic
 - No new tasks

KDG Runtime

Application Programmer

KDG Specializations

Opportunities

- Each step needs a barrier
- Task's rw-set potentially updated several times prior to execution

Specializations

- Remove unnecessary barriers given program properties
- Construct KDG for prefix of tasks
- Use different DAG representations

Program Properties

- Sources
 - **In general**: safe source test could inspect entire state
 - Local safe source test
 - Stable source (sources are safe)
- RW-sets
 - **In general**: task could change other tasks' RW-sets
 - Non-increasing RW-set
 - Structure-based RW-set
- New tasks
 - **In general**: task could create new tasks at any priority
 - Monotonic
 - No new tasks

KDG Runtime

Application Programmer

KDG Specializations

Opportunities

- Each step needs a barrier
- Task's rw-set potentially updated several times prior to execution

Specializations

- Remove unnecessary barriers given program properties
- Construct KDG for prefix of tasks
- Use different DAG representations

KDG Runtime

Program Properties

- Sources
 - In general: safe source test could inspect entire state
 - Local safe source test
 - Stable source (sources are safe)
- RW-sets
 - In general: task could change other tasks' RW-sets
 - Non-increasing RW-set
 - Structure-based RW-set
- New tasks
 - In general: task could create new tasks at any priority
 - Monotonic
 - No new tasks

Application Programmer

KDGs in Action


- Application programmer
 - Writes programs with ordered loops
 - Loop body must use provided data structure **library**
 - Loop properties provided via **annotations**
 - Loop body must read all elements before modifying any (**cautious**)
- KDG runtime (in Galois system)
 - Uses library API to track rw-sets at runtime
 - Selects executor based on annotations


```
Graph g = ...
Set<Event> E = ...


@hasStructuredRWSets
@monotonic
foreach Event e in E orderedby e.t
    process(e, g)
    if *
        E.push(newE)
```

Evaluation

- 7 applications
 - From billiards simulation (**hard-to-parallelize**) to tree traversal (**well supported** by prior work)
- 2 types of programs
 - **Other**
 - 3rd-party handwritten, application-specific, OpenMP tasks, Cilk
 - **KDG-Auto**
 - Ordered loops + program property annotations
- 40-core, shared-memory machine







Summary

- **Problem**
 - Dependence graph scheduling is insufficient for many ordered programs
- **Solution**
 - Develop general KDG executor
 - Specialize executor according to small number of program properties

“It is a mistake to try to look too far ahead. The chain of destiny can only be grasped one link at a time.”

-Winston Churchill

Kinetic Dependence Graphs

M. Amber Hassaan,

Donald Nguyen,

Keshav Pingali

The University of Texas at Austin

