

Deterministic Galois: On-demand, Portable and Parameterless


Donald Nguyen

Andrew Lenharth, Keshav Pingali

The University of Texas at Austin

Why Determinism?

- Parallelism introduces non-determinism from scheduling



- Goal:** Eliminate scheduling non-determinism
 - Simplify debugging, fault tolerance

Desired Qualities

- **On-demand**
 - Determinism can be expensive; allow users to easily enable determinism as desired
- **Portable**
 - Deterministic result should be the same regardless of machine architecture, including number of threads
- **Parameterless**
 - There should be no user-tunable parameters that affect output

Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]

Avenues to Determinism

Determinism by Construction


- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]
- On-demand?
- Portable?
- Parameterless?

Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]

- On-demand?
- Portable?
- Parameterless?



Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]

• On-demand?



• Portable?



• Parameterless?

Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]

• On-demand?



• Portable?



• Parameterless?



Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]


• On-demand?



• Portable?



• Parameterless?



Determinism by Scheduling


- Provide deterministic version of **low-level** scheduling primitives
 - Determinism w/o rewriting programs
- Examples
 - Kendo [Olszewski09]
 - RCDC [Deviatti11]

Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]

- On-demand?
- Portable?
- Parameterless?



Determinism by Scheduling

- Provide deterministic version of **low-level** scheduling primitives
 - Determinism w/o rewriting programs
- Examples
 - Kendo [Olszewski09]
 - RCDC [Deviatti11]

- On-demand?
- Portable?
- Parameterless?

Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]

• On-demand?



• Portable?



• Parameterless?



Determinism by Scheduling

- Provide deterministic version of **low-level** scheduling primitives
 - Determinism w/o rewriting programs
- Examples
 - Kendo [Olszewski09]
 - RCDC [Devietti11]

• On-demand?



• Portable?

• Parameterless?

Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]

• On-demand?



• Portable?



• Parameterless?



Determinism by Scheduling

- Provide deterministic version of **low-level** scheduling primitives
 - Determinism w/o rewriting programs
- Examples
 - Kendo [Olszewski09]
 - RCDC [Deviatti11]

• On-demand?



• Portable?



• Parameterless?

Avenues to Determinism

Determinism by Construction

- Programs must conform to set of deterministic program constructs
- Examples
 - Fork-join [Blumofe95]
 - DPJ [Bocchino11]
 - PBBS [Blelloch11]


• On-demand?



• Portable?



• Parameterless?




Determinism by Scheduling

- Provide deterministic version of **low-level** scheduling primitives
 - Determinism w/o rewriting programs
- Examples
 - Kendo [Olszewski09]
 - RCDC [Deviatti11]

• On-demand?



• Portable?



• Parameterless?



Avenues to Determinism

High-level Non-deterministic Programming Model

- Deterministic Galois (this work)
- Non-determinism expressed in constructs beyond threads and locks
 - System responsible for deterministic execution if desired
- Targeted towards fine-grain tasks
 - 10--1000 cycles per task
 - Frequent communication

Avenues to Determinism

High-level Non-deterministic Programming Model

- Deterministic Galois (this work)
- Non-determinism expressed in constructs beyond threads and locks
 - System responsible for deterministic execution if desired
- Targeted towards fine-grain tasks
 - 10--1000 cycles per task
 - Frequent communication

- On-demand?
- Portable?
- Parameterless?

Avenues to Determinism

High-level Non-deterministic Programming Model

- Deterministic Galois (this work)
- Non-determinism expressed in constructs beyond threads and locks
 - System responsible for deterministic execution if desired
- Targeted towards fine-grain tasks
 - 10--1000 cycles per task
 - Frequent communication

- On-demand?
- Portable?
- Parameterless?



Avenues to Determinism

High-level Non-deterministic Programming Model

- Deterministic Galois (this work)
- Non-determinism expressed in constructs beyond threads and locks
 - System responsible for deterministic execution if desired
- Targeted towards fine-grain tasks
 - 10--1000 cycles per task
 - Frequent communication

- On-demand?
- Portable?
- Parameterless?



Avenues to Determinism

High-level Non-deterministic Programming Model

- Deterministic Galois (this work)
- Non-determinism expressed in constructs beyond threads and locks
 - System responsible for deterministic execution if desired
- Targeted towards fine-grain tasks
 - 10--1000 cycles per task
 - Frequent communication


- On-demand?
- Portable?
- Parameterless?



Galois Programming Model

- **Galois system:** runtime and library of concurrent data structures
- **Set iterators** express parallelism
 - **Operator:** function to apply
 - **Neighborhood:** data accessed
 - **Cautious:** operators read their entire neighborhood before writing to any element
 - **Failsafe point:** Point between reading and writing
- Set iterator produces **some serialization** of operator invocations
 - Only source of non-determinism in Galois programs


```
Graph g
for_each (Edge e : wl)
  Node n = g.getEdgeDst(e)
  ...
  for (Node m :
    g.out_edges(n))
    ... g.getData(m) ...
  ...
```



Galois Programming Model

- **Galois system:** runtime and library of concurrent data structures
- **Set iterators** express parallelism
 - **Operator:** function to apply
 - **Neighborhood:** data accessed
 - **Cautious:** operators read their entire neighborhood before writing to any element
 - **Failsafe point:** Point between reading and writing
- Set iterator produces **some serialization** of operator invocations
 - Only source of non-determinism in Galois programs


```
Graph g
for_each (Edge e : wl)
  Node n = g.getEdgeDst(e)
  ...
  for (Node m :
    g.out_edges(n))
    ... g.getData(m) ...
  ...
```



Galois Programming Model

- **Galois system:** runtime and library of concurrent data structures
- **Set iterators** express parallelism
 - **Operator:** function to apply
 - **Neighborhood:** data accessed
 - **Cautious:** operators read their entire neighborhood before writing to any element
 - **Failsafe point:** Point between reading and writing
- Set iterator produces **some serialization** of operator invocations
 - Only source of non-determinism in Galois programs


```
Graph g
for_each (Edge e : wl)
  Node n = g.getEdgeDst(e)
  ...
  for (Node m :
    g.out_edges(n))
    ... g.getData(m) ...
  ...
```



Galois Programming Model


- **Galois system:** runtime and library of concurrent data structures
- **Set iterators** express parallelism
 - **Operator:** function to apply
 - **Neighborhood:** data accessed
 - **Cautious:** operators read their entire neighborhood before writing to any element
 - **Failsafe point:** Point between reading and writing
- Set iterator produces **some serialization** of operator invocations
 - Only source of non-determinism in Galois programs

```
Graph g
for_each (Edge e : wl)
  Node n = g.getEdgeDst(e)
  ...
  for (Node m :
    g.out_edges(n))
    ... g.getData(m) ...
  ...
```




Deterministic Interference Graph Scheduling

- **Traditional** Galois execution
 - Asynchronous, non-deterministic
- **Deterministic** Galois execution
 - Construct an **interference graph**
 - Execute independent set
 - Repeat




Deterministic Interference Graph Scheduling

- **Traditional** Galois execution
 - Asynchronous, non-deterministic
- **Deterministic** Galois execution
 - Construct an **interference graph**
 - Execute independent set
 - Repeat




Deterministic Interference Graph Scheduling

- **Traditional** Galois execution
 - Asynchronous, non-deterministic
- **Deterministic** Galois execution
 - Construct an **interference graph**
 - Execute independent set
 - Repeat




Deterministic Interference Graph Scheduling

- **Traditional** Galois execution
 - Asynchronous, non-deterministic
- **Deterministic** Galois execution
 - Construct an **interference graph**
 - Execute independent set
 - Repeat




Deterministic Interference Graph Scheduling

- **Traditional** Galois execution
 - Asynchronous, non-deterministic
- **Deterministic** Galois execution
 - Construct an **interference graph**
 - Execute independent set
 - Repeat




Deterministic Interference Graph Scheduling

- **Traditional** Galois execution
 - Asynchronous, non-deterministic
- **Deterministic** Galois execution
 - Construct an **interference graph**
 - Execute independent set
 - Repeat




Deterministic Interference Graph Scheduling

- **Traditional** Galois execution
 - Asynchronous, non-deterministic
- **Deterministic** Galois execution
 - Construct an **interference graph**
 - Execute independent set
 - Repeat



DIG Scheduling


- Deterministically select independent set
 - Form total order on tasks
 - Select tasks that are least among direct neighbors



$A < B < C < D$

DIG Scheduling


- Deterministically select independent set
 - Form total order on tasks
 - Select tasks that are least among direct neighbors



$A < B < C < D$

DIG Scheduling in Practice


- **Implicitly** build neighborhoods and interference graph
- **Marks** associated with data structure elements
- Acquire marks by writing task id atomically
 - Overwriting an id only if replacing greater value
- Execute tasks whose neighborhood only contains their own marks
- Final mark values are **deterministic**



$$A < B < C$$

DIG Scheduling in Practice


- **Implicitly** build neighborhoods and interference graph
- **Marks** associated with data structure elements
- Acquire marks by writing task id atomically
 - Overwriting an id only if replacing greater value
- Execute tasks whose neighborhood only contains their own marks
- Final mark values are **deterministic**



$$A < B < C$$

DIG Scheduling in Practice


- **Implicitly** build neighborhoods and interference graph
- **Marks** associated with data structure elements
- Acquire marks by writing task id atomically
 - Overwriting an id only if replacing greater value
- Execute tasks whose neighborhood only contains their own marks
- Final mark values are **deterministic**



$$A < B < C$$

DIG Scheduling in Practice


- **Implicitly** build neighborhoods and interference graph
- **Marks** associated with data structure elements
- Acquire marks by writing task id atomically
 - Overwriting an id only if replacing greater value
- Execute tasks whose neighborhood only contains their own marks
- Final mark values are **deterministic**



$A < B < C$

DIG Scheduling in Practice


- **Implicitly** build neighborhoods and interference graph
- **Marks** associated with data structure elements
- Acquire marks by writing task id atomically
 - Overwriting an id only if replacing greater value
- Execute tasks whose neighborhood only contains their own marks
- Final mark values are **deterministic**



$A < B < C$

DIG Scheduling in Practice


- **Implicitly** build neighborhoods and interference graph
- **Marks** associated with data structure elements
- Acquire marks by writing task id atomically
 - Overwriting an id only if replacing greater value
- Execute tasks whose neighborhood only contains their own marks
- Final mark values are **deterministic**



$$A < B < C$$

DIG Scheduling in Practice

- **Implicitly** build neighborhoods and interference graph
- **Marks** associated with data structure elements
- Acquire marks by writing task id atomically
 - Overwriting an id only if replacing greater value
- Execute tasks whose neighborhood only contains their own marks
- Final mark values are **deterministic**




DIG Scheduling in Practice

- Sequence of rounds
 - Round has two phases
- **Phase 1:** Inspect neighborhoods
 - Execute operator to its **failsafe point**
 - Acquire marks
- **Phase 2:** Execute roots
 - Reexecute operator, checking mark values
 - Postpone any task that did not read its marks

```
for_each (Edge e : wl)
  Node n = g.getEdgeDst(e)
  ...
  for (Node m :
    g.out_edges(n))
    g.getData(m).value += 1
```

...

Failsafe Point




DIG Scheduling in Practice

- Sequence of rounds
 - Round has two phases
- **Phase 1:** Inspect neighborhoods
 - Execute operator to its **failsafe point**
 - Acquire marks
- **Phase 2:** Execute roots
 - Reexecute operator, checking mark values
 - Postpone any task that did not read its marks

```
for_each (Edge e : wl)
  Node n = g.getEdgeDst(e)
  ...
  for (Node m :
    g.out_edges(n))
    g.getData(m).value += 1
```

...

Failsafe Point



Optimizations

- Resuming tasks
 - Avoid reexecuting tasks
 - Suspend and resume execution at failsafe point
 - Provide buffers to save local state
- Windowing
 - Inspect only a subset of tasks at a time
 - Adaptive algorithm varies window size

Evaluation

Platform

- Intel 4x10 core machine


Applications

- PBBS [Blelloch11]
 - Breadth-first search (**bfs**)
 - Delaunay mesh refinement (**dmr**)
 - Delaunay triangulation (**dt**)
 - Maximal independent set (**mis**)

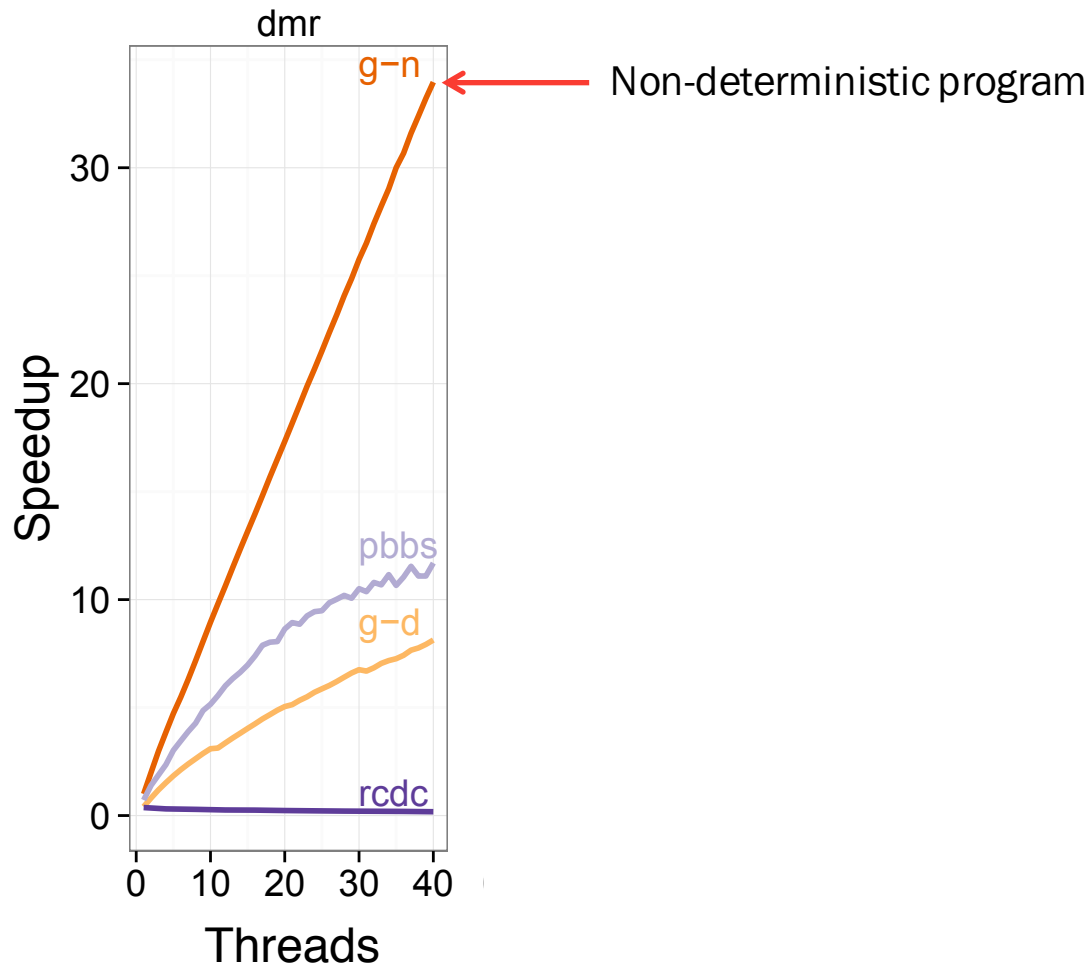
Deterministic Systems

- RCDC
 - Deterministic by scheduling
 - Implementation of Kendo algorithm
- PBBS
 - Deterministic by construction
 - Handwritten deterministic implementations
- **Deterministic Galois**
 - Non-deterministic programs (**g-n**) automatically made deterministic (**g-d**)

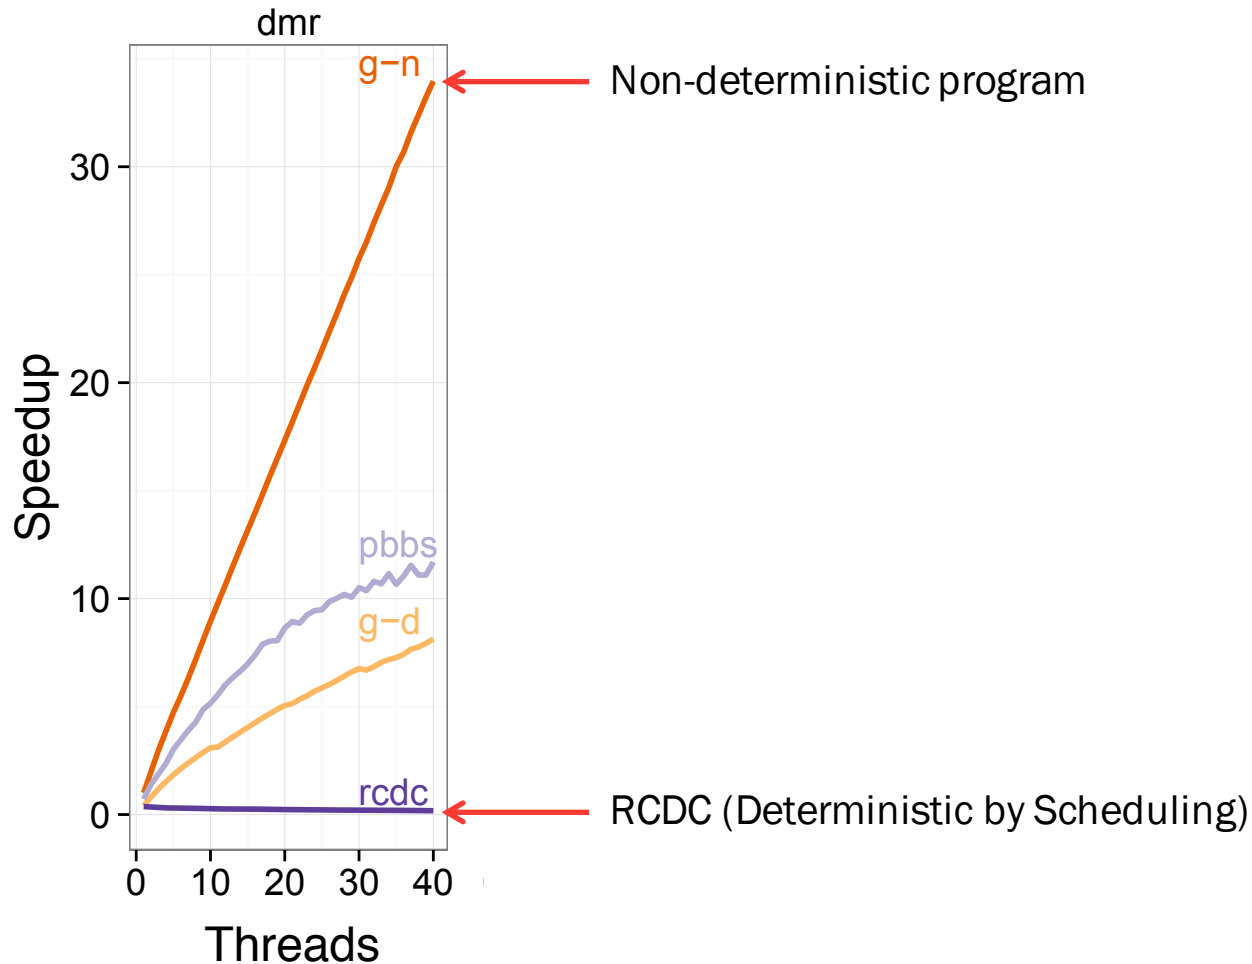
Evaluation



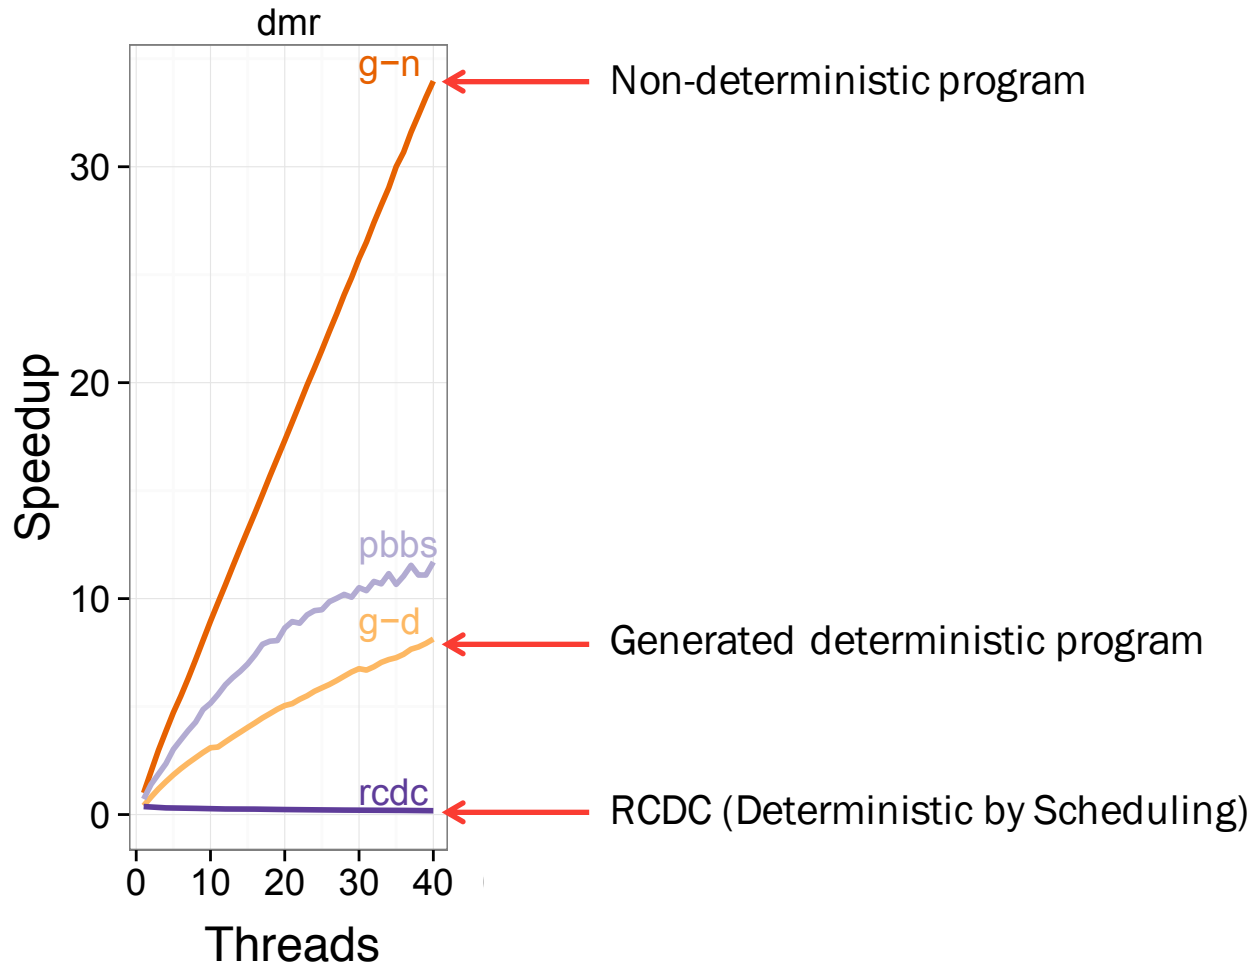
Evaluation




Evaluation




Evaluation



Evaluation




Evaluation



Sources of Overhead

- Additional instructions
- Extending critical path by round execution
- Deterministic schedule chosen may not be optimal
- **Locality**
 - Inspecting a task and executing it are separated in time

Locality



Conclusion

- Make programs deterministic using deterministic interference graph scheduling
 - Facilitates **on-demand**, **portable** and **parameterless** deterministic programs
 - Generated deterministic programs comparable to handwritten ones
 - Non-deterministic programs often much faster than deterministic ones
- In the paper
 - More applications, machines
 - Quantify locality under different systems
 - Measure impact of optimizations

<http://iss.ices.utexas.edu/galois>

Deterministic Galois: On-demand, Portable and Parameterless

Donald Nguyen

Andrew Lenharth, Keshav Pingali

The University of Texas at Austin