

# **Synthesizing Concurrent Schedulers for Irregular Algorithms**

**Donald Nguyen** and Keshav Pingali


University of Texas at Austin

# The Scheduling Problem



- Given a decomposition of program into tasks or *activities*, determine
  - a) the assignment of activities to processors **and**
  - b) the order in which each processor executes its activities
- In this work, focus on dynamic scheduling


# The Scheduling Problem



- Given a decomposition of program into tasks or *activities*, determine
  - a) the assignment of activities to processors **and**
  - b) the order in which each processor executes its activities
- In this work, focus on dynamic scheduling


# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation



# Scheduling Is Important


- Find the shortest path from source to all nodes in graph via iterative relaxation






# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation



# Scheduling Is Important


- Find the shortest path from source to all nodes in graph via iterative relaxation



- Policies
  - Distance Order, FIFO, D-stepping

# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




- Policies
  - Distance Order, FIFO, D-stepping

	t=1
Dist. order	0.5 s

# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




- Policies
  - Distance Order, FIFO, D-stepping

	t=1
Dist. order	0.5 s
FIFO	> 2.5 hours

# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




- Policies
  - Distance Order, FIFO, D-stepping

	t=1	t=4
Dist. order	0.5 s	20 s
FIFO	> 2.5 hours	~ 1 hour

# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation




- Policies
  - Distance Order, FIFO, D-stepping

	t=1	t=4
Dist. order	0.5 s	20 s
FIFO	> 2.5 hours	~ 1 hour
D-stepping	0.5 s	0.4 s

# Scheduling Is Important

- Find the shortest path from source to all nodes in graph via iterative relaxation



- Policies
  - Distance Order, FIFO, D-stepping

	t=1	t=4
Dist. order	0.5 s	20 s
FIFO	> 2.5 hours	~ 1 hour
D-stepping	0.5 s	0.4 s

**Order matters even in unordered algorithms**



# Scheduling Is Hard

## Sequential Scheduling

```
while (item = wl.poll())  
    fn(item)  
    wl.add(item)
```

# Scheduling Is Hard

## Sequential Scheduling

```
while (item = wl.poll())  
    fn(item)  
    wl.add(item)
```

**Tedious when exploring new policies**

# Scheduling Is Hard

## Sequential Scheduling

```
while (item = wl.poll())  
  fn(item)  
  wl.add(item)
```

## Parallel Scheduling

```
foreach (item : S)  
  fn(item)
```

```
spawn A, B, C  
sync
```

**Tedious when exploring new policies**

# Scheduling Is Hard

## Sequential Scheduling

```
while (item = wl.poll())  
  fn(item)  
  wl.add(item)
```

Tedious when exploring new policies

## Parallel Scheduling

```
foreach (item : S)  
  fn(item)
```

```
spawn A, B, C  
sync
```

At most,  $k$  scheduling policies ( $k \geq 2$ )

# Scheduling Is Hard

## Sequential Scheduling

```
while (item = wl.poll())  
  fn(item)  
  wl.add(item)
```

Tedious when exploring new policies

## Parallel Scheduling

```
foreach (item : S)  
  fn(item)
```

```
spawn A, B, C  
sync
```

OpenMP ( $k = 0$ )  
.Net TPL ( $k = 1$ )  
TBB ( $k = 2$ )

At most,  $k$  scheduling policies ( $k \leq 2$ )

# Scheduling Is Hard

## Sequential Scheduling

```
while (item = wl.poll())  
  fn(item)  
  wl.add(item)
```

Tedious when exploring new policies

## Parallel Scheduling

```
foreach (item : S)  
  fn(item)
```

```
spawn A, B, C  
sync
```

OpenMP ( $k = 0$ )  
.Net TPL ( $k = 1$ )  
TBB ( $k = 2$ )

At most,  $k$  scheduling policies ( $k \leq 2$ )

New policies *same as* making  
concurrent, scalable data structures

# Scheduling Is Hard

## Sequential Scheduling

```
while (item = wl.poll())  
  fn(item)  
  wl.add(item)
```

Tedious when exploring new policies

## Parallel Scheduling

```
foreach (item : S)  
  fn(item)
```

```
spawn A, B, C  
sync
```

OpenMP ( $k = 0$ )  
.Net TPL ( $k = 1$ )  
TBB ( $k = 2$ )

At most,  $k$  scheduling policies ( $k \leq 2$ )

New policies *same as* making  
concurrent, scalable data structures

```
GaloisRuntime.foreach(S, fn, SchedulingPolicy);
```

# Contributions

- A language for scheduling policies
  - *Declarative*: sophisticated schedulers w/o writing code
  - *Effective*: performance comparable to hand-written and often better than previous schedulers

Get good performance without writing  
(serial or concurrent) scheduling code




# Outline

1. Language
2. Synthesis
3. Results

$P$	$::=$	Global: $D$ Local: $D$	<i>Specification</i>
$D$	$::=$	$R_{NF}^* \quad R_F^?$	<i>Ordering</i>
$R_{NF}$	$::=$	ChunkedFIFO( $k$ )	<i>Non-final rule</i>
		ChunkedLIFO( $k$ )	
		Ordered( $f : x, y \rightarrow \text{bool}$ )	
		OrderedByMetric( $f : x \rightarrow \mathbb{R}$ )	
$R_F$	$::=$	FIFO	<i>Final rule</i>
		LIFO	
		Random	
$k$	$::=$	...	<i>Integer</i>


# Rules and their composition



order



# Rules and their composition





OrderedByMetric( $g$ )

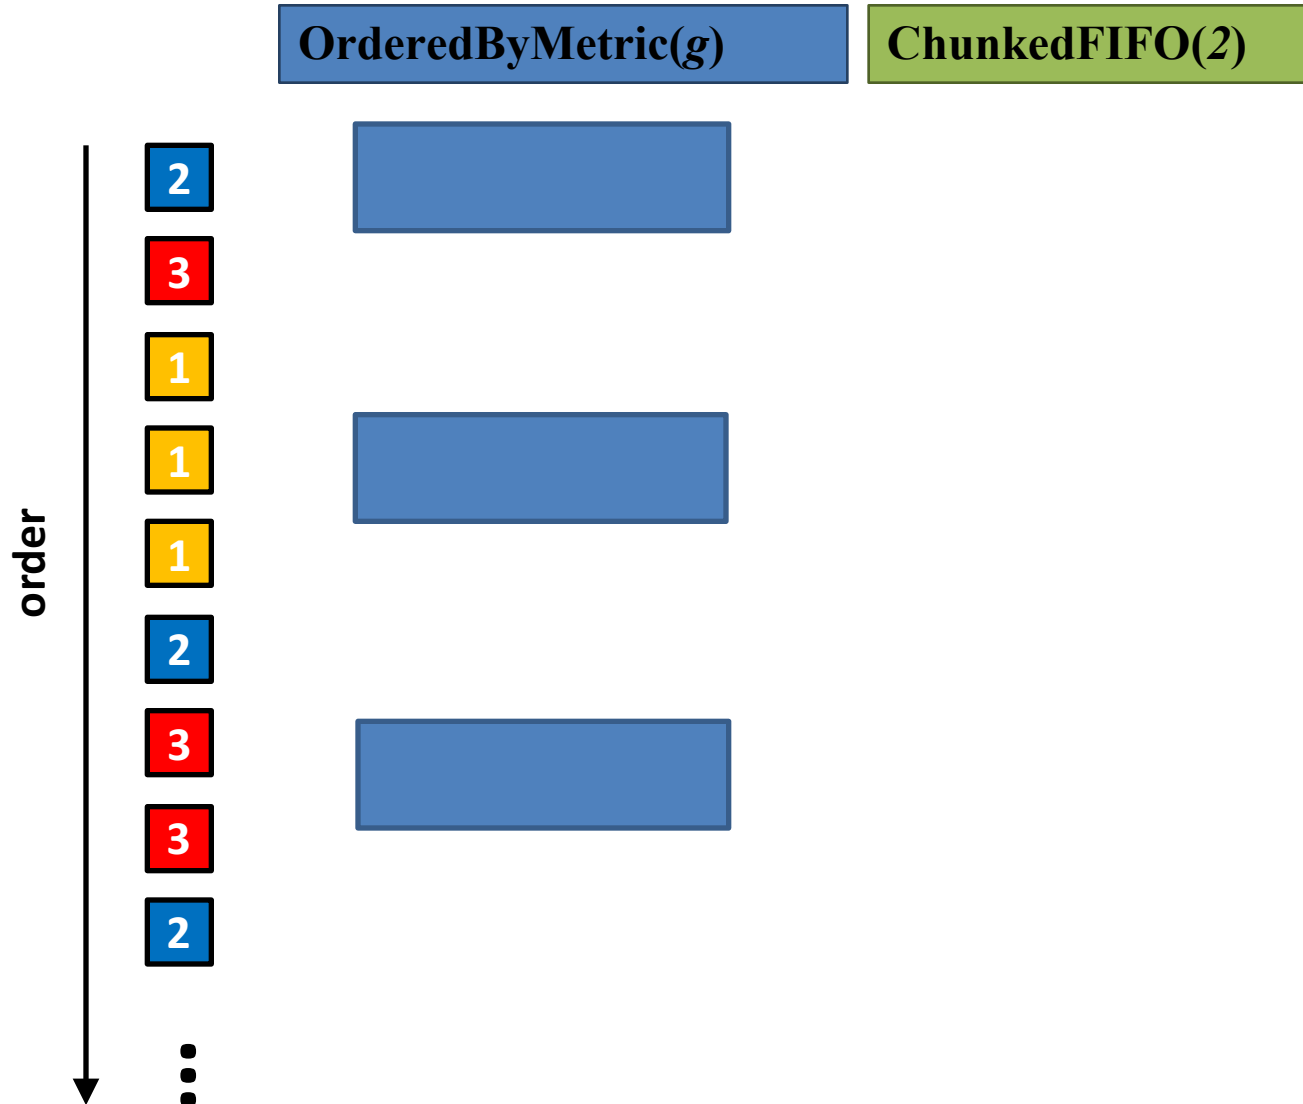

order



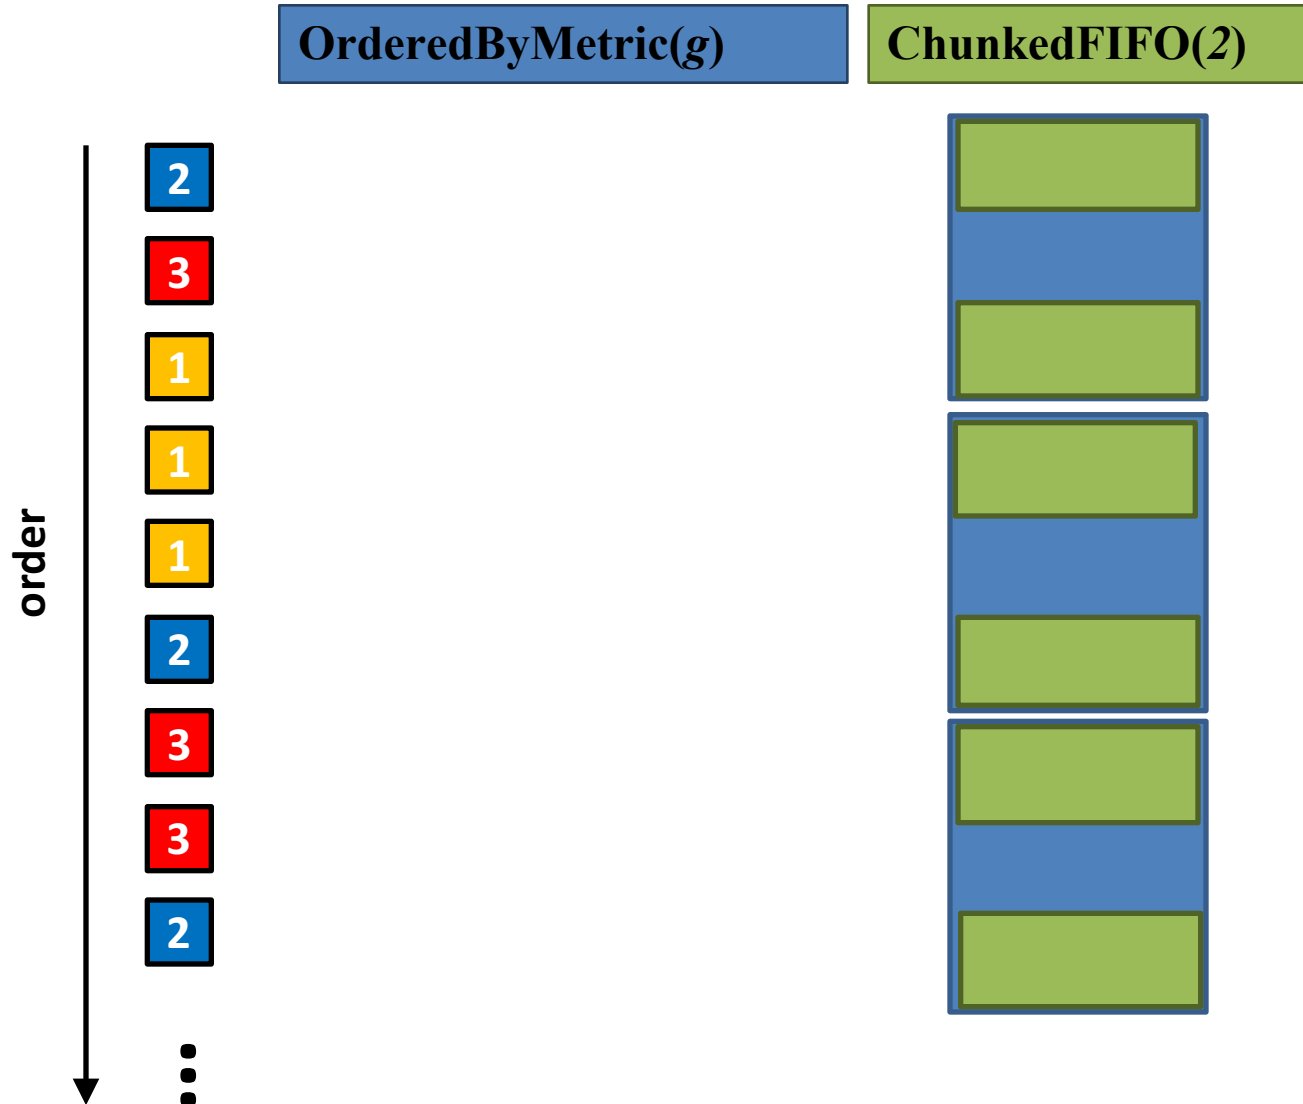

# Rules and their composition





# Rules and their composition





# Rules and their composition



# Rules and their composition



# Rules and their composition





# Application-specific Policies

App	Order	Scheduling Policy	
PFP	FIFO	FIFO	[Goldberg88]
PFP	HL	OrderedByMetric( <del><math>n</math></del> . $-n.height$ ) FIFO	[Cherkassy95]

# Application-specific Policies

App	Order	Scheduling Policy	
PFP	FIFO	FIFO	[Goldberg88]
PFP	HL	OrderedByMetric( $\ast n. -n.height$ ) FIFO	[Cherkassy95]
SSSP	D-stepping	OrderedByMetric( $\ast n. \frac{n.w}{D} + \dots$ ) FIFO	[Meyer98]
SSSP	Dijkstra	Ordered( $\ast a, b. a.w \bullet b.w$ )	[Dijkstra59]

# Application-specific Policies

App	Order	Scheduling Policy	
PFP	FIFO	FIFO	[Goldberg88]
PFP	HL	OrderedByMetric( $\star n. -n.height$ ) FIFO	[Cherkassy95]
SSSP	D-stepping	OrderedByMetric( $\star n. \frac{n.w}{D} + \dots$ ) FIFO	[Meyer98]
SSSP	Dijkstra	Ordered( $\star a, b. a.w \bullet b.w$ )	[Dijkstra59]
DMR	Local stack	ChunkedFIFO( $k$ ) Local: LIFO	[Kulkarni08]

# Application-specific Policies

App	Order	Scheduling Policy	
PFP	FIFO	FIFO	[Goldberg88]
PFP	HL	OrderedByMetric( $\star n. -n.height$ ) FIFO	[Cherkassy95]
SSSP	D-stepping	OrderedByMetric( $\star n. \frac{n.w}{D} + \dots$ ) FIFO	[Meyer98]
SSSP	Dijkstra	Ordered( $\star a, b. a.w \dots b.w$ )	[Dijkstra59]
DMR	Local stack	ChunkedFIFO( $k$ ) Local: LIFO	[Kulkarni08]
DT	BRIO	OrderedByMetric( $\star p. p.rnd$ ) ChunkedFIFO( $k$ )	[Amenta03]



# Application-specific Policies

App	Order	Scheduling Policy	
PFP	FIFO	FIFO	[Goldberg88]
PFP	HL	OrderedByMetric( $\star n. -n.height$ ) FIFO	[Cherkassy95]
SSSP	D-stepping	OrderedByMetric( $\star n. \frac{n.w}{D} + \dots$ ) FIFO	[Meyer98]
SSSP	Dijkstra	Ordered( $\star a, b. a.w \bullet b.w$ )	[Dijkstra59]
DMR	Local stack	ChunkedFIFO( $k$ ) Local: LIFO	[Kulkarni08]
DT	BRIO	OrderedByMetric( $\star p. p.rnd$ ) ChunkedFIFO( $k$ )	[Amenta03]
PTA	Split		[Nielson99]

# Outline

1. Language
- 2. Synthesis**
3. Results

$P$	$::=$	Global: $D$ Local: $D$	<i>Specification</i>
$D$	$::=$	$R_{NF}^* R_F?$	<i>Ordering</i>
$R_{NF}$	$::=$	ChunkedFIFO( $k$ )	<i>Non-final rule</i>
		ChunkedLIFO( $k$ )	
		Ordered( $f : x, y \rightarrow \text{bool}$ )	
		OrderedByMetric( $f : x \rightarrow \mathbb{R}$ )	
$R_F$	$::=$	FIFO	<i>Final rule</i>
		LIFO	
		Random	
$k$	$::=$	...	<i>Integer</i>



# Synthesis of Serial Schedulers

**OrderedByMetric( $g$ )**

**ChunkedFIFO( $k$ )**



...

# Synthesis of Serial Schedulers

**OrderedByMetric( $g$ )**

**ChunkedFIFO( $k$ )**

...

**void add(T t)**

**T poll()**



# Synthesis of Serial Schedulers

**OrderedByMetric( $g$ )**

**ChunkedFIFO( $k$ )**

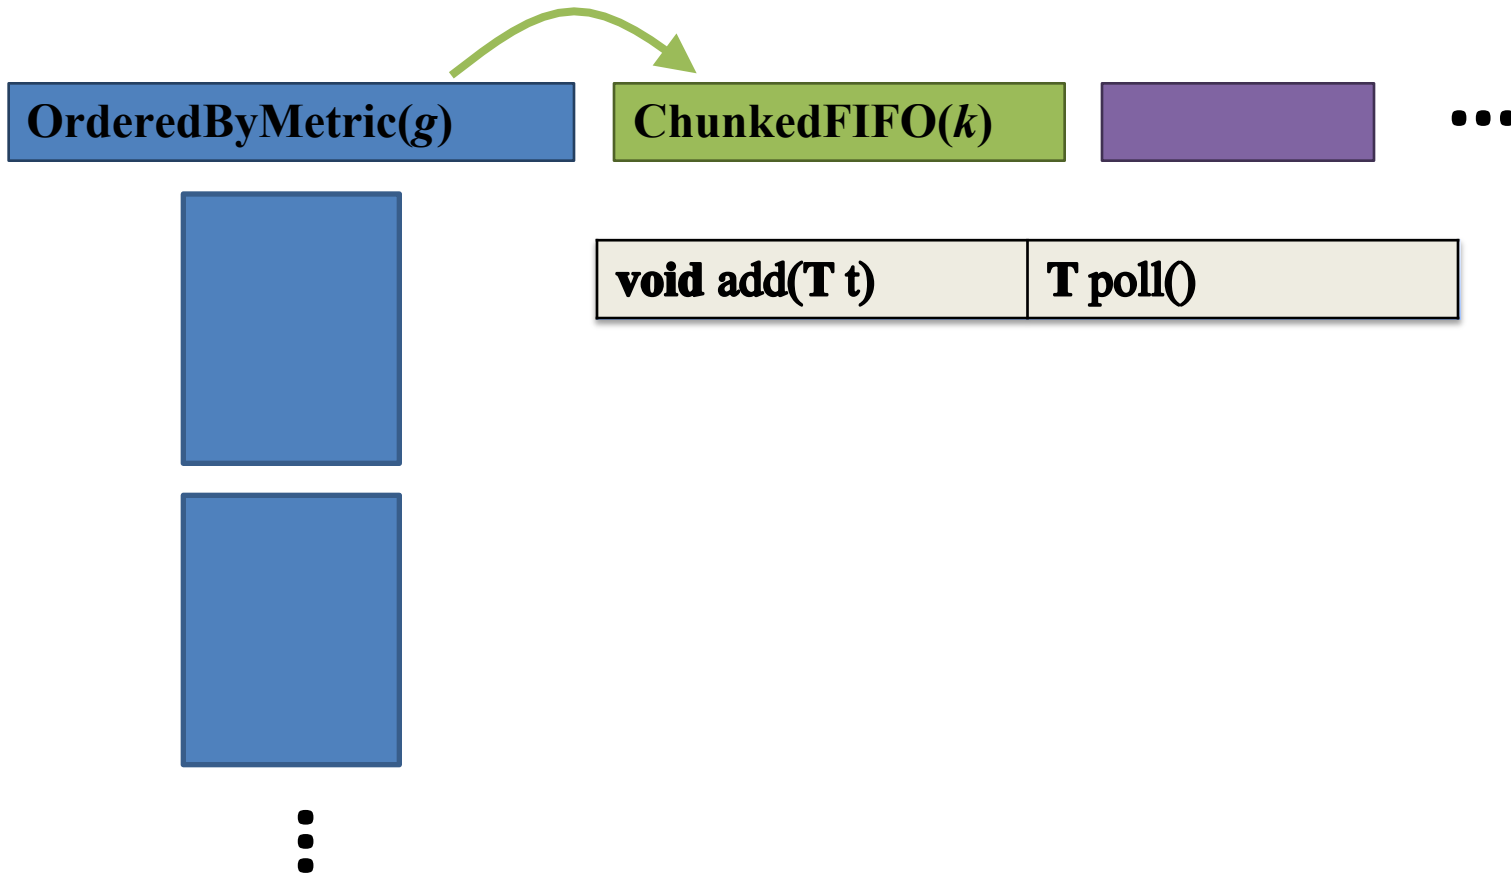
...

**void add(T t)**

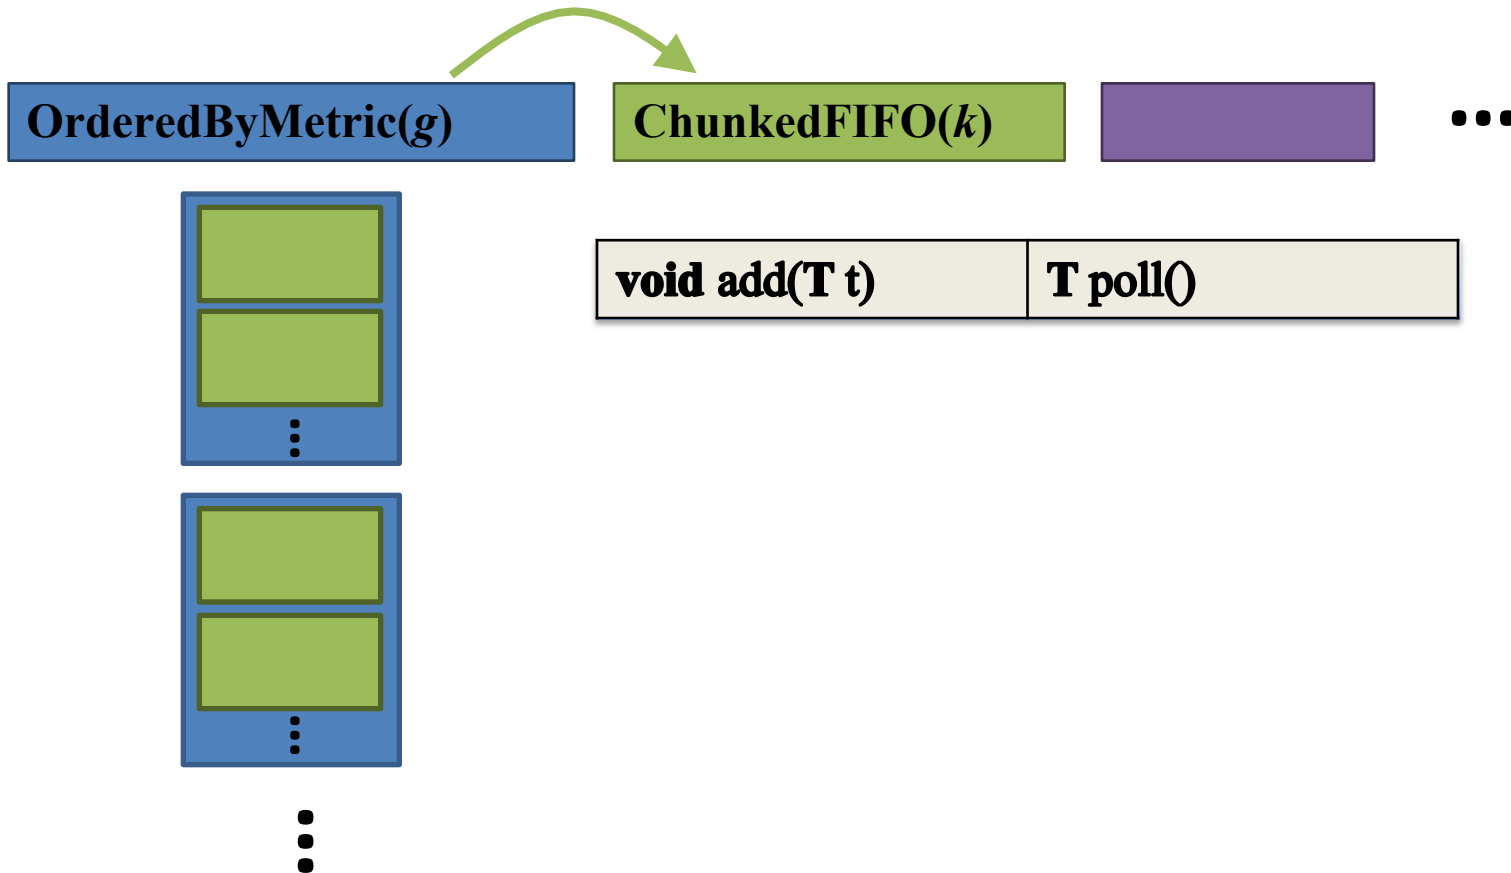
**T poll()**

⋮

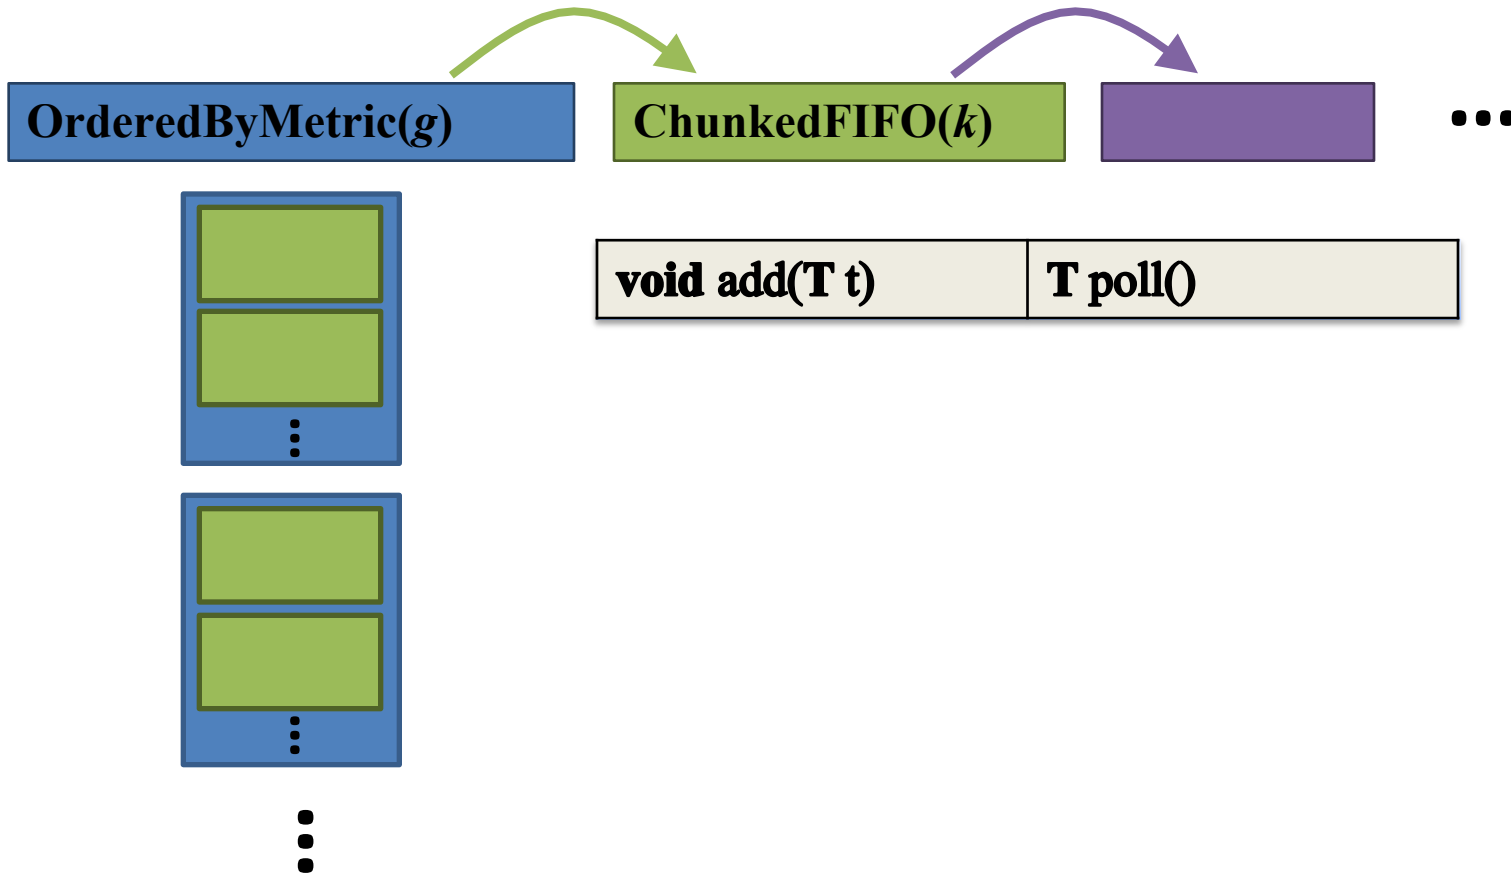
# Synthesis of Serial Schedulers



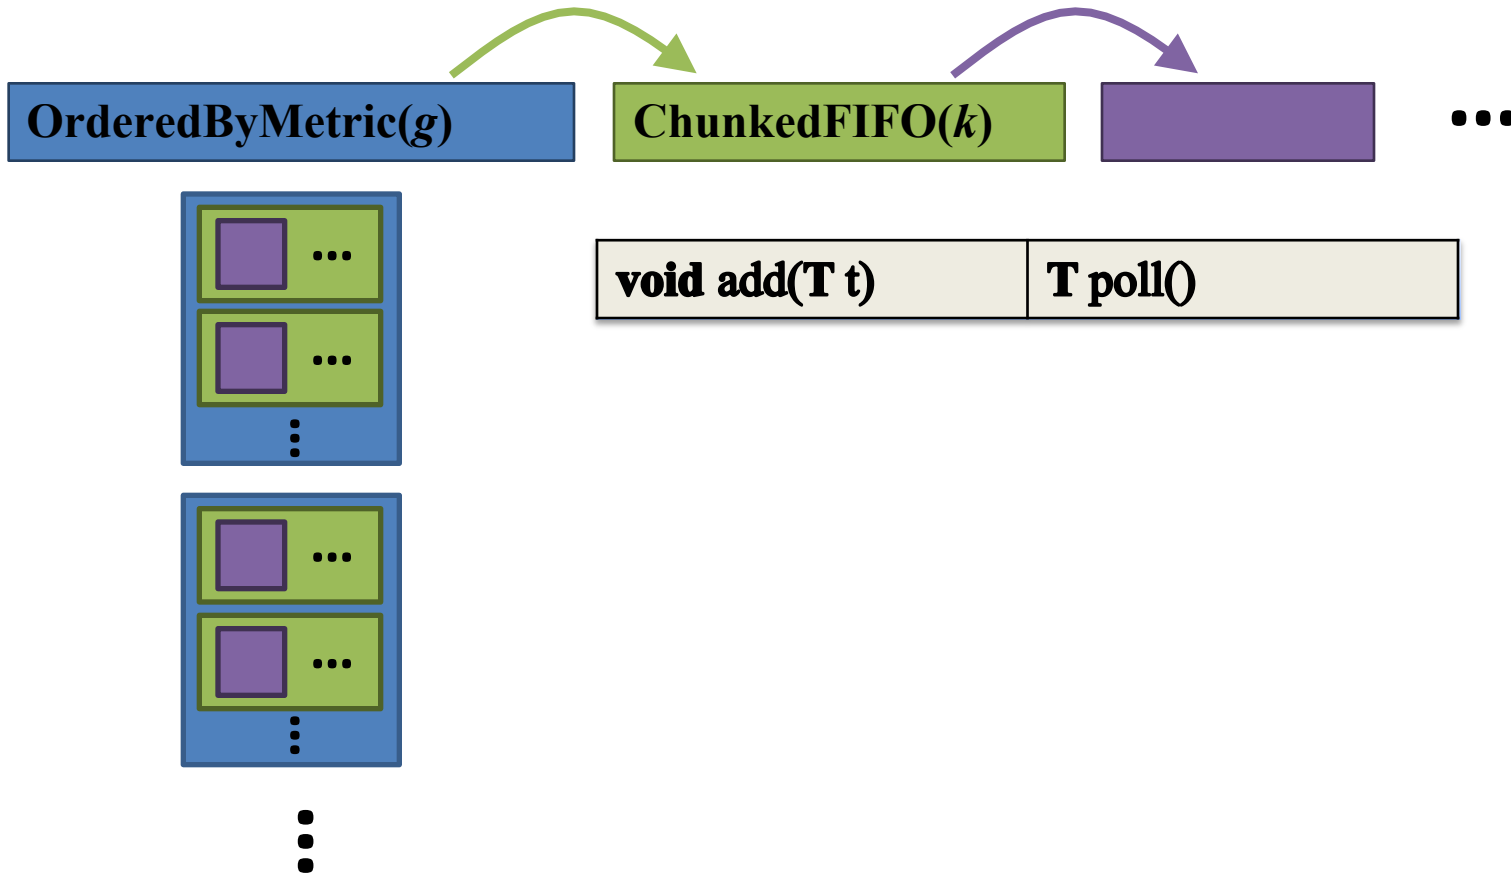
# Synthesis of Serial Schedulers





# Synthesis of Serial Schedulers



# Synthesis of Serial Schedulers



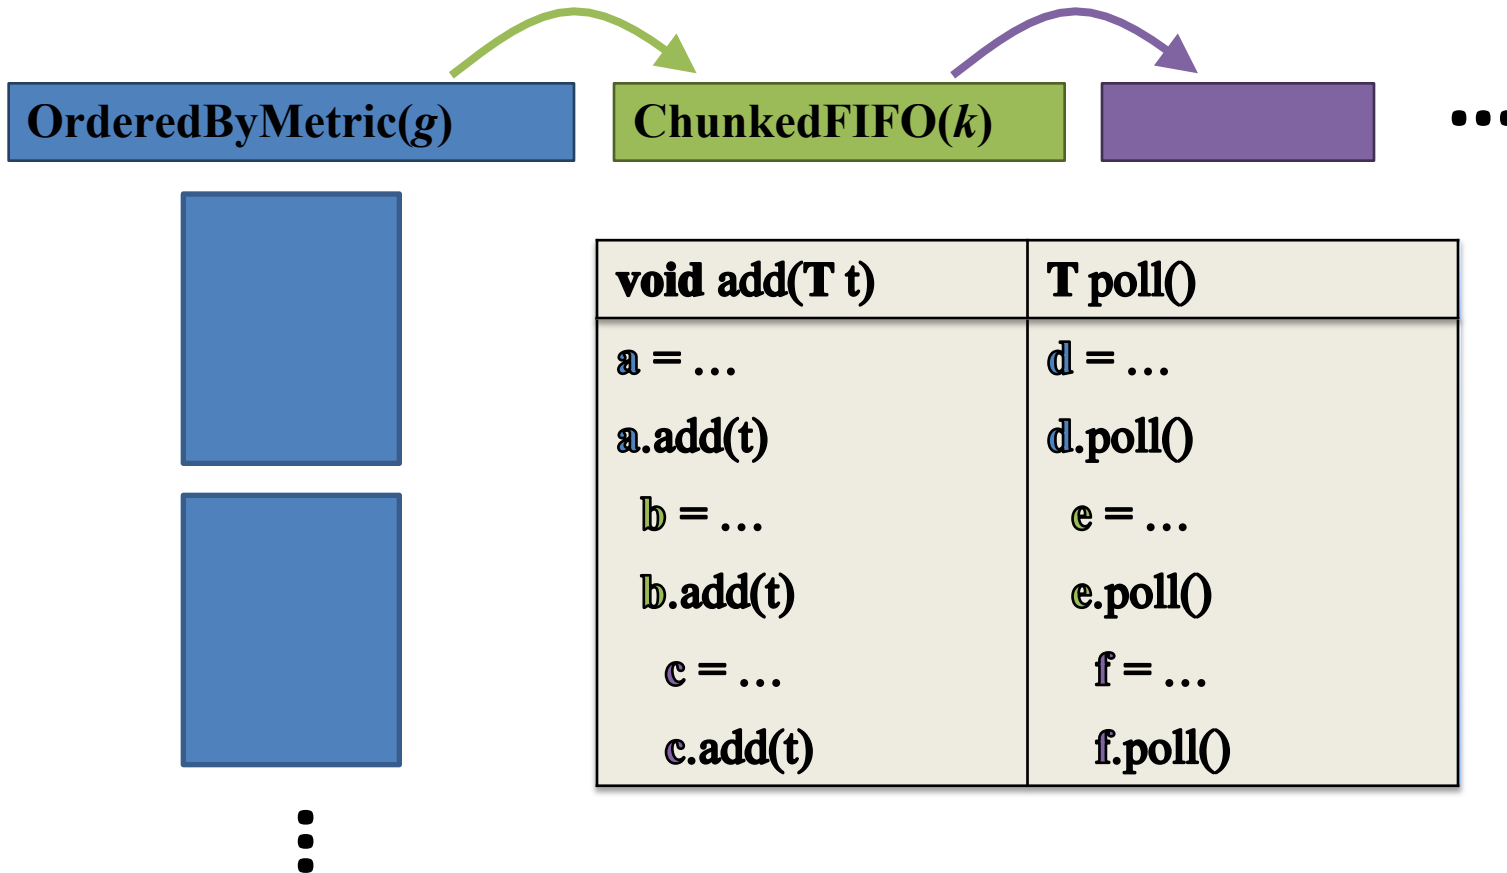
# Synthesis of Serial Schedulers



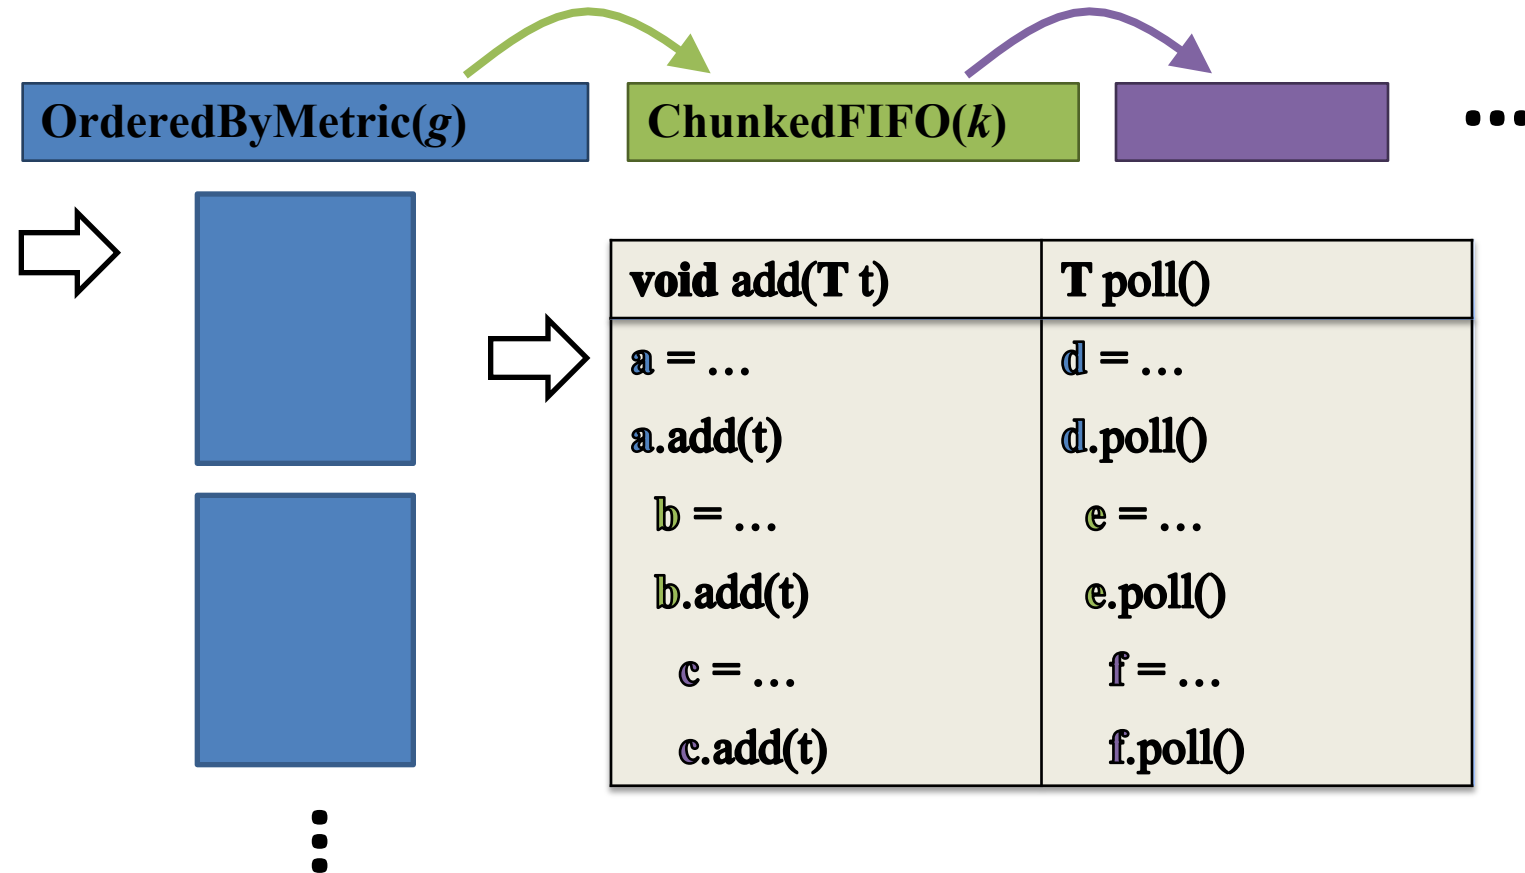
⋮

<b>void add(T t)</b>	<b>T poll()</b>
<b>a</b> = ...	<b>d</b> = ...
<b>a.add(t)</b>	<b>d.poll()</b>
<b>b</b> = ...	<b>e</b> = ...
<b>b.add(t)</b>	<b>e.poll()</b>
<b>c</b> = ...	<b>f</b> = ...
<b>c.add(t)</b>	<b>f.poll()</b>

# Synthesis of Serial Schedulers

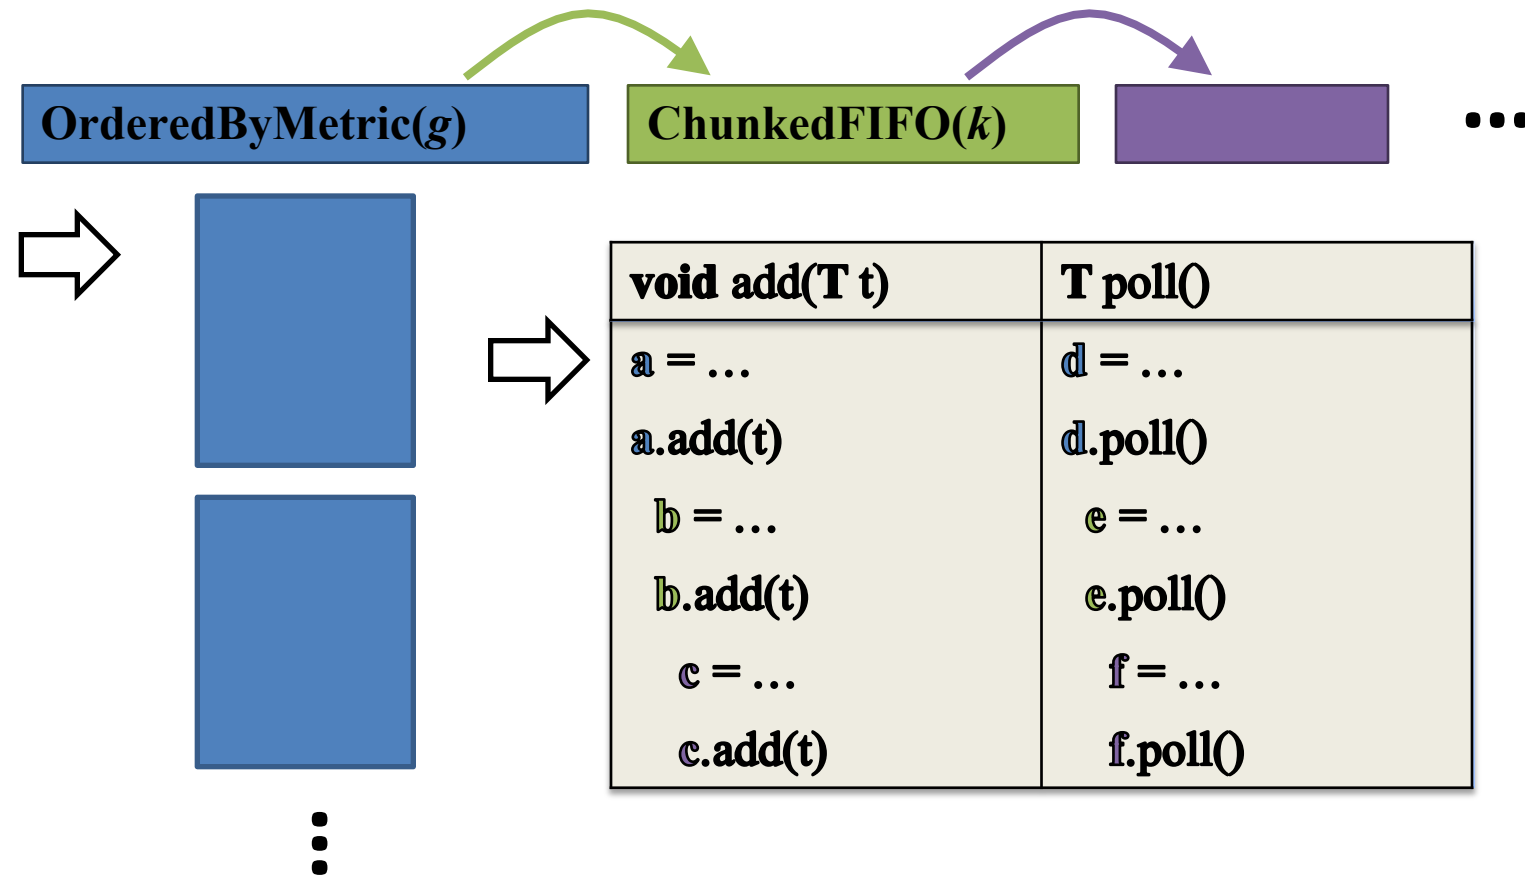


# Synthesis of Serial Schedulers

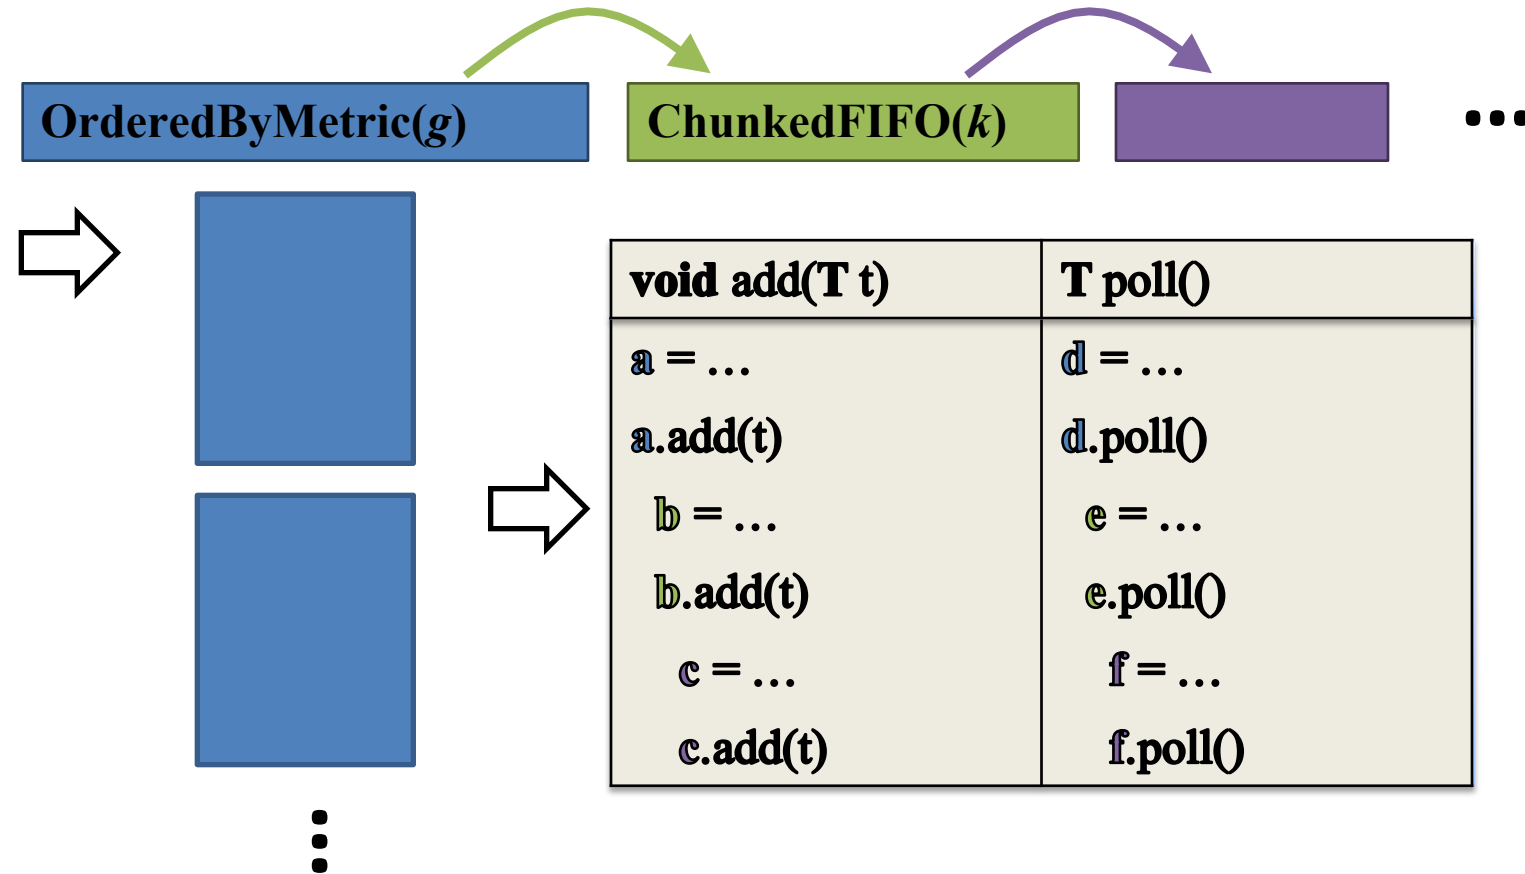





# Synthesis of Serial Schedulers



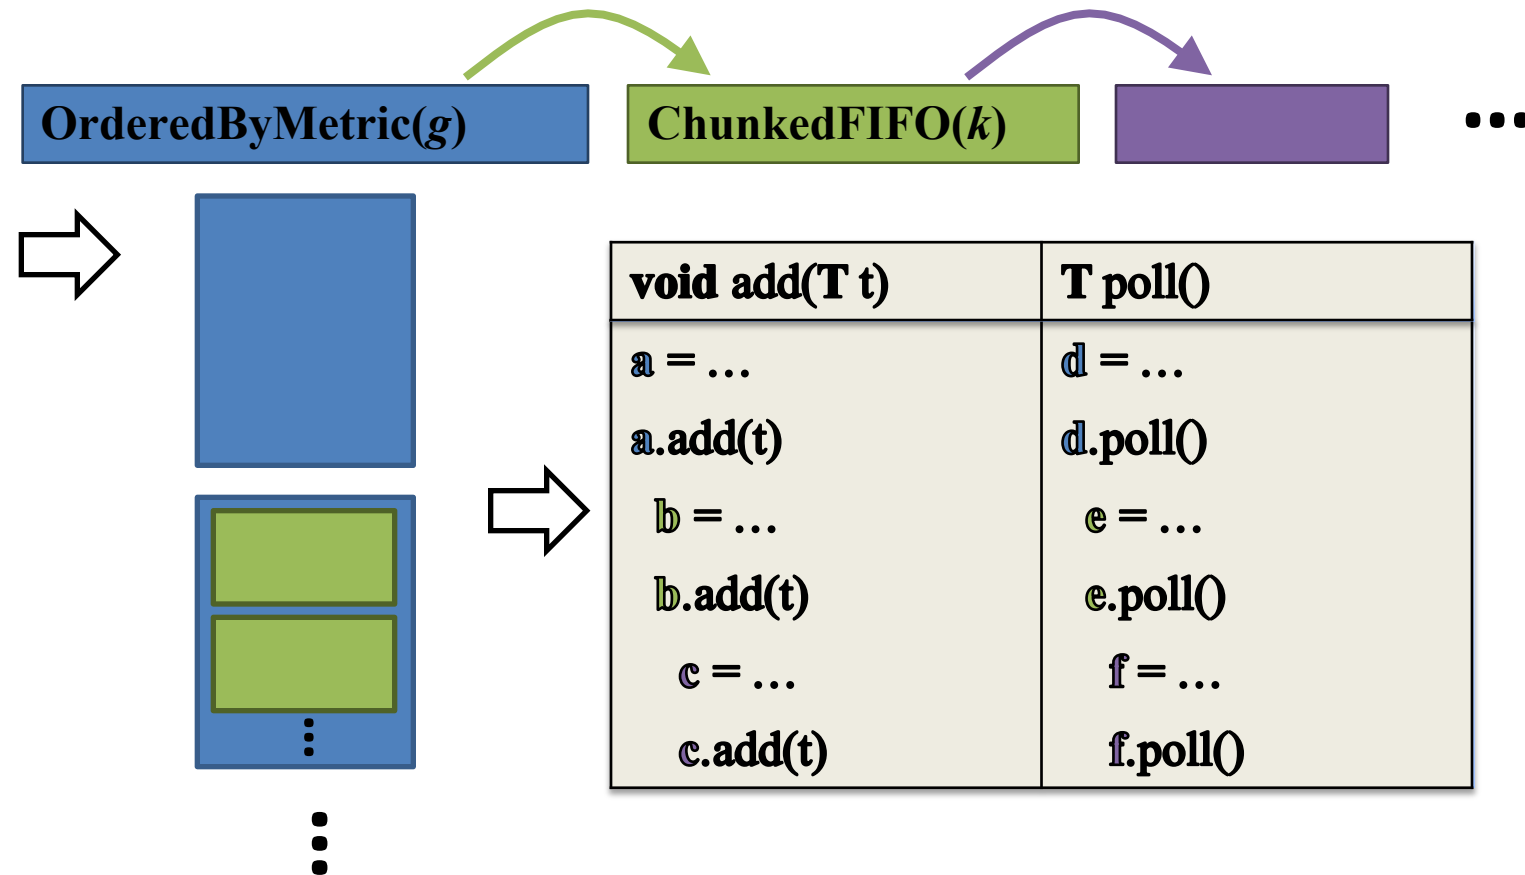
# Synthesis of Serial Schedulers



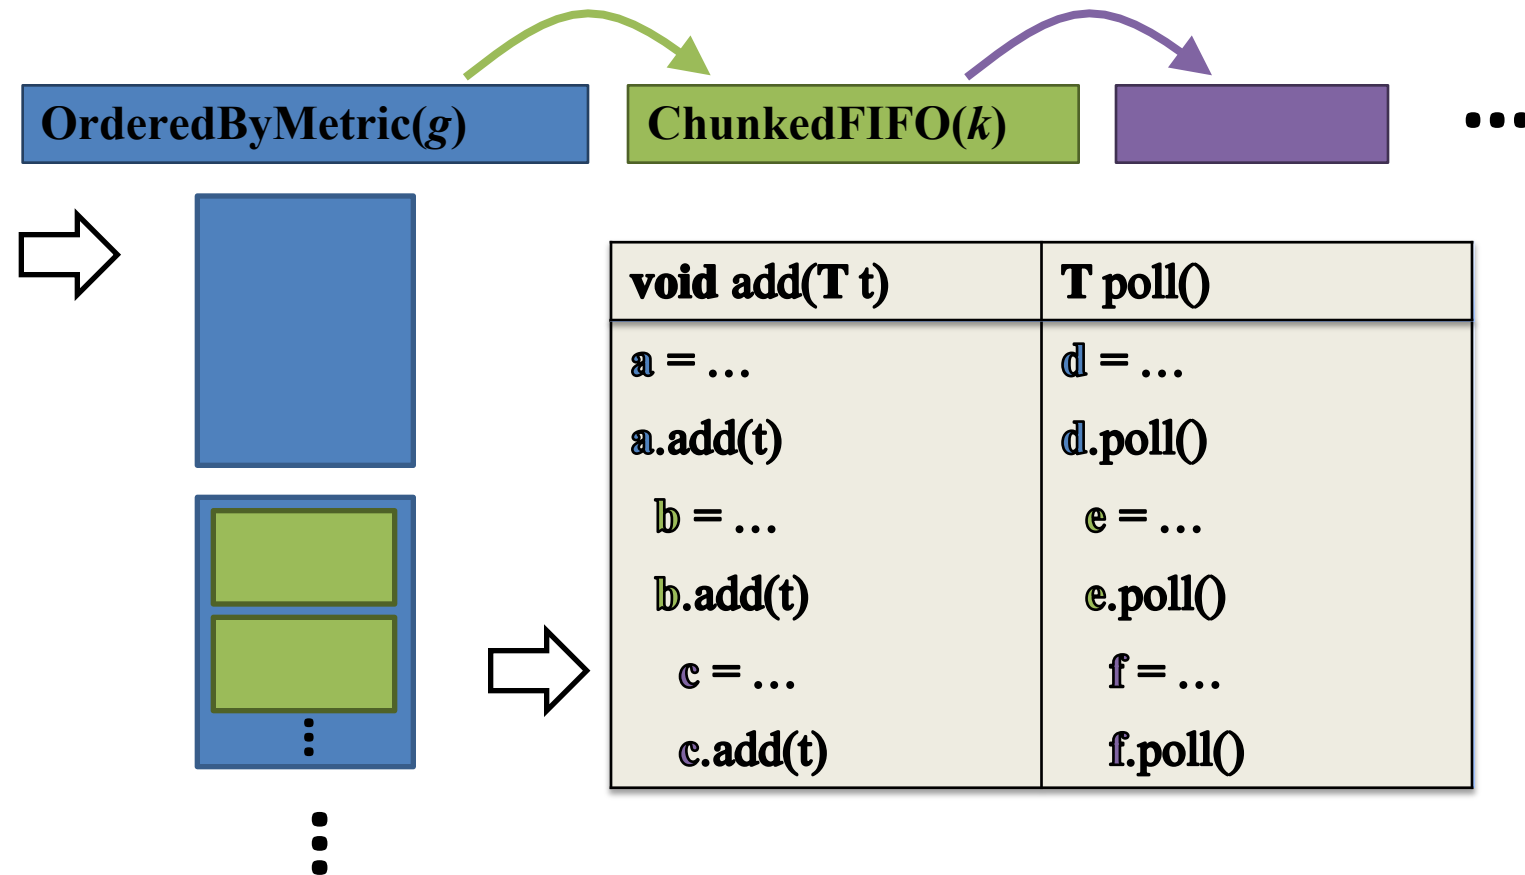
# Synthesis of Serial Schedulers



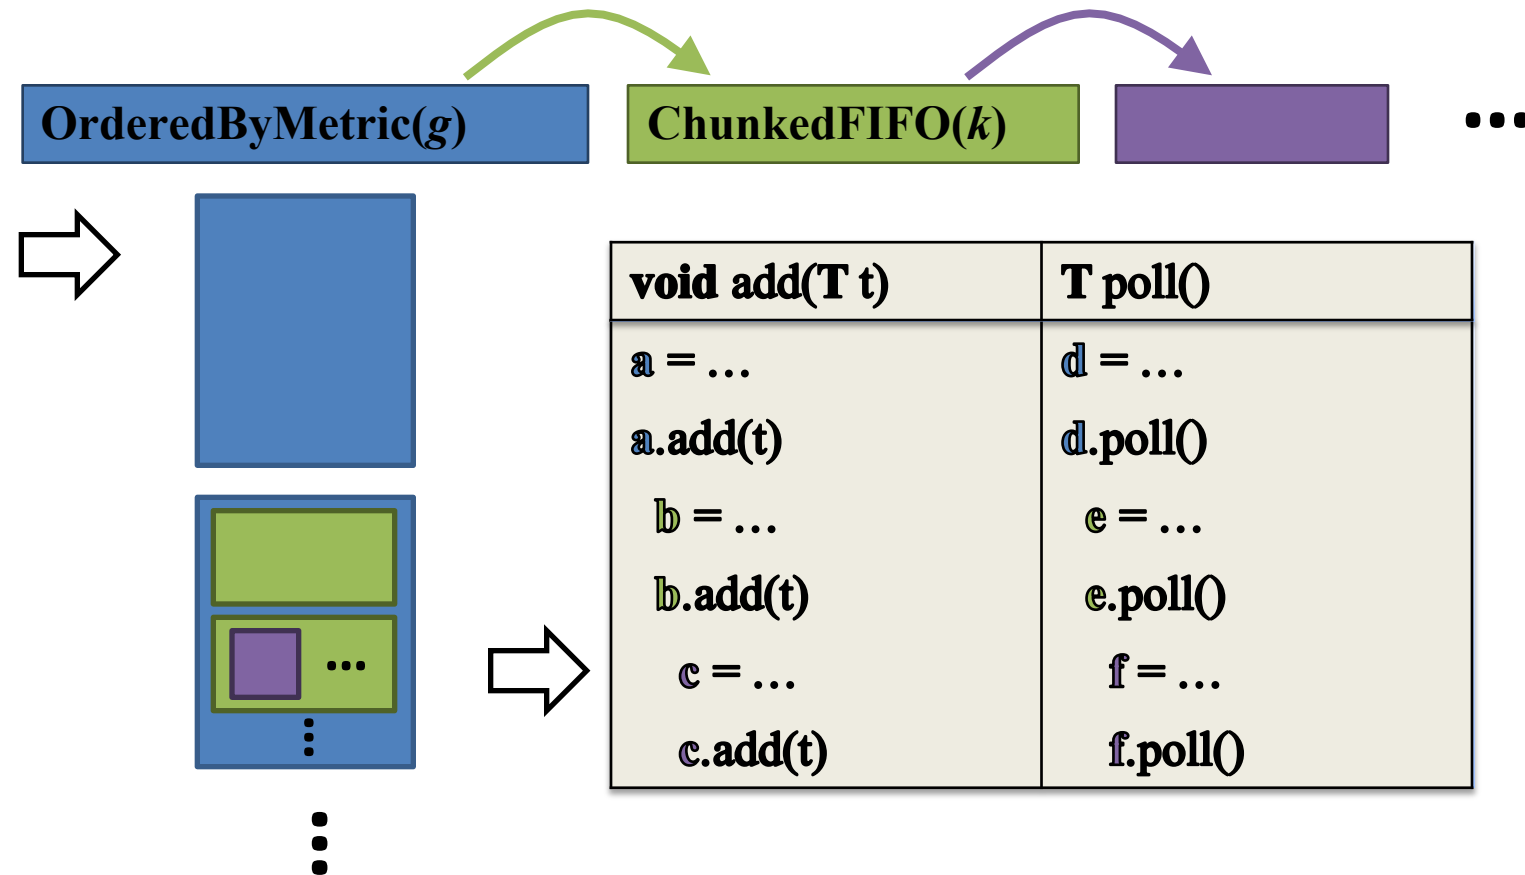
# Synthesis of Serial Schedulers



# Synthesis of Serial Schedulers



# Synthesis of Serial Schedulers



# Synthesis of Concurrent Schedulers

<b>void add(T t)</b>	<b>T poll()</b>
<b>a</b> = ... <b>a.add(t)</b> <b>b</b> = ... <b>b.add(t)</b> <b>c</b> = ... <b>c.add(t)</b>	<b>d</b> = ... <b>d.poll()</b> <b>e</b> = ... <b>e.poll()</b> <b>f</b> = ... <b>f.poll()</b>


# Synthesis of Concurrent Schedulers

atomic


<b>void add(T t)</b>	<b>T poll()</b>
<div><div>a = ...</div><div>a.add(t)</div></div> <div><div>b = ...</div><div>b.add(t)</div></div> <div><div>c = ...</div><div>c.add(t)</div></div>	<div><div>d = ...</div><div>d.poll()</div></div> <div><div>e = ...</div><div>e.poll()</div></div> <div><div>f = ...</div><div>f.poll()</div></div>



# Synthesis of Concurrent Schedulers



# Synthesis of Concurrent Schedulers



- Key property
  - Scheduling orders are heuristics
  - **Scheduler is valid as long as we eventually retrieve every item added (in some order)**

# Synthesis of Concurrent Schedulers

atomic

<code>void add(T t)</code>	<code>T poll()</code>
<code>a = ...</code> <code>a.add(t)</code>	<code>d = ...</code> <code>d.poll()</code>
<code>b = ...</code> <code>b.add(t)</code>	<code>e = ...</code> <code>e.poll()</code>
<code>c = ...</code> <code>c.add(t)</code>	<code>f = ...</code> <code>f.poll()</code>

- Key property
  - Scheduling orders are heuristics
  - **Scheduler is valid as long as we eventually retrieve every item added (in some order)**


# Synthesis of Concurrent Schedulers

	<b>void add(T t)</b>	<b>T poll()</b>	<b>T poll-s()</b>
atomic	<b>a</b> = ... <b>a.add(t)</b>	<b>d</b> = ... <b>d.poll()</b>	<b>d</b> = ... <b>d.poll()</b>
	<b>b</b> = ... <b>b.add(t)</b>	<b>e</b> = ... <b>e.poll()</b>	<b>e</b> = ... <b>e.poll()</b>
	<b>c</b> = ... <b>c.add(t)</b>	<b>f</b> = ... <b>f.poll()</b>	<b>f</b> = ... <b>f.poll()</b>

- Key property
  - Scheduling orders are heuristics
  - **Scheduler is valid as long as we eventually retrieve every item added (in some order)**

# Outline

1. Language
2. Synthesis
- 3. Results**



# Methodology

- Three machines
  - Nehalem: 2x4-core @ 2.93GHz (Xeon X5570)
  - Shanghai: 4x4-core @ 2.7GHz (Opteron 8384)
  - Niagara: 4x8-core @ 1.4GHz (UltraSPARC T2 Plus)
- Java
  - Sun JDK 1.6
  - 20GB heap
  - Last of three runs within same VM instance

# Scheduling Policies


- Fixed-forms
  - **Rand**
  - **LIFO, FIFO**: Global queue or stack
  - **WS-L, WS-F**: Work-stealing with queue or stack
  - **BS-L, BS-F**
- Synthesized
  - **Base**: ChunkedFIFO(32)
  - **AS**: Application-specific policy

# Application-specific Policies


App	Order	Scheduling Policy	
PFP	FIFO	FIFO	[Goldberg88]
PFP	HL	OrderedByMetric( $\ast n. -n.height$ ) FIFO	[Cherkassy95]
SSSP	D-stepping	OrderedByMetric( $\ast n. \frac{n.w}{D} + \dots$ ) FIFO	[Meyer98]
SSSP	Dijkstra	Ordered( $\ast a, b. a.w \leq b.w$ )	[Dijkstra59]
DMR	Local stack	ChunkedFIFO( $k$ ) Local: LIFO	[Kulkarni08]
DT	BRIO	OrderedByMetric( $\ast p. p.rnd$ ) ChunkedFIFO( $k$ )	[Amenta03]
PTA	Split	$BS-F$	[Nielson99]




# Serial Runtime: DMR



# Serial Runtime




# Speedup: DMR




# Speedup


nehalem



shanghai




niagara

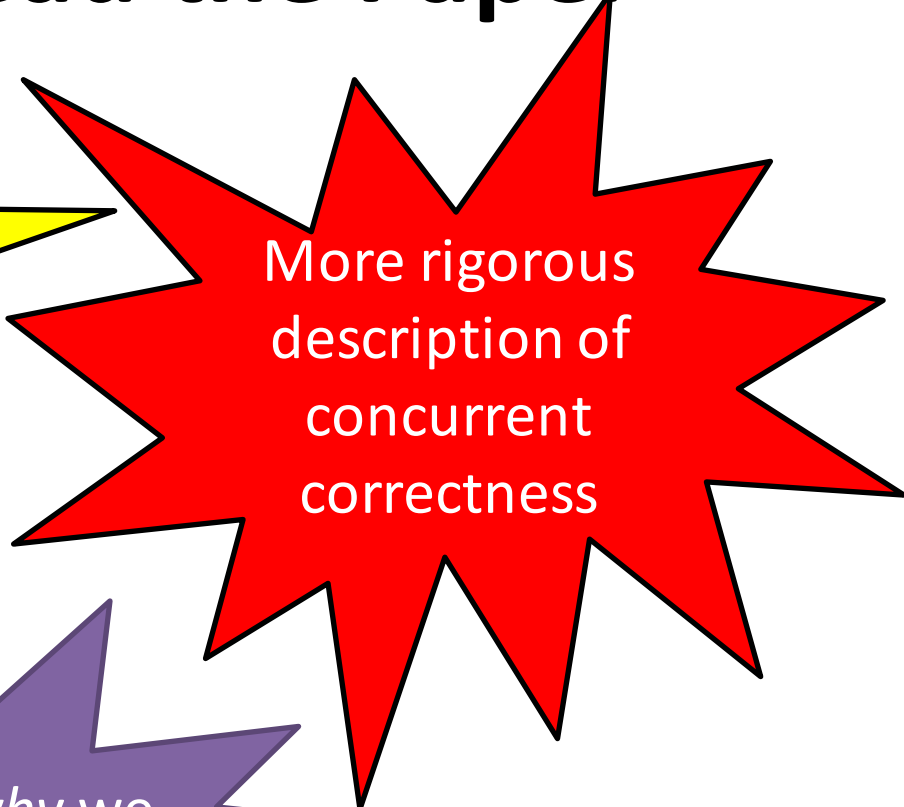


AS versus WS-F


# An Advert to Read the Paper

A yellow starburst shape with a black outline, containing text.

**Optimizing  
implementation  
of synthesized  
worklists**

A red starburst shape with a black outline, containing text.

More rigorous  
description of  
concurrent  
correctness

A purple starburst shape with a black outline, containing text.

Insight on *why* we  
see the reported  
speedups

# Conclusion

1. High-level programmer control over scheduling is important for serial ***and*** parallel performance
  2. Providing such control is easy and extensible
- Future work
    - Architectural optimizations
    - More algorithms

<http://iss.ices.utexas.edu/galois>