

What Scalable Programs Need from Transactional Memory

Donald Nguyen

d.nguyen@synthace.com

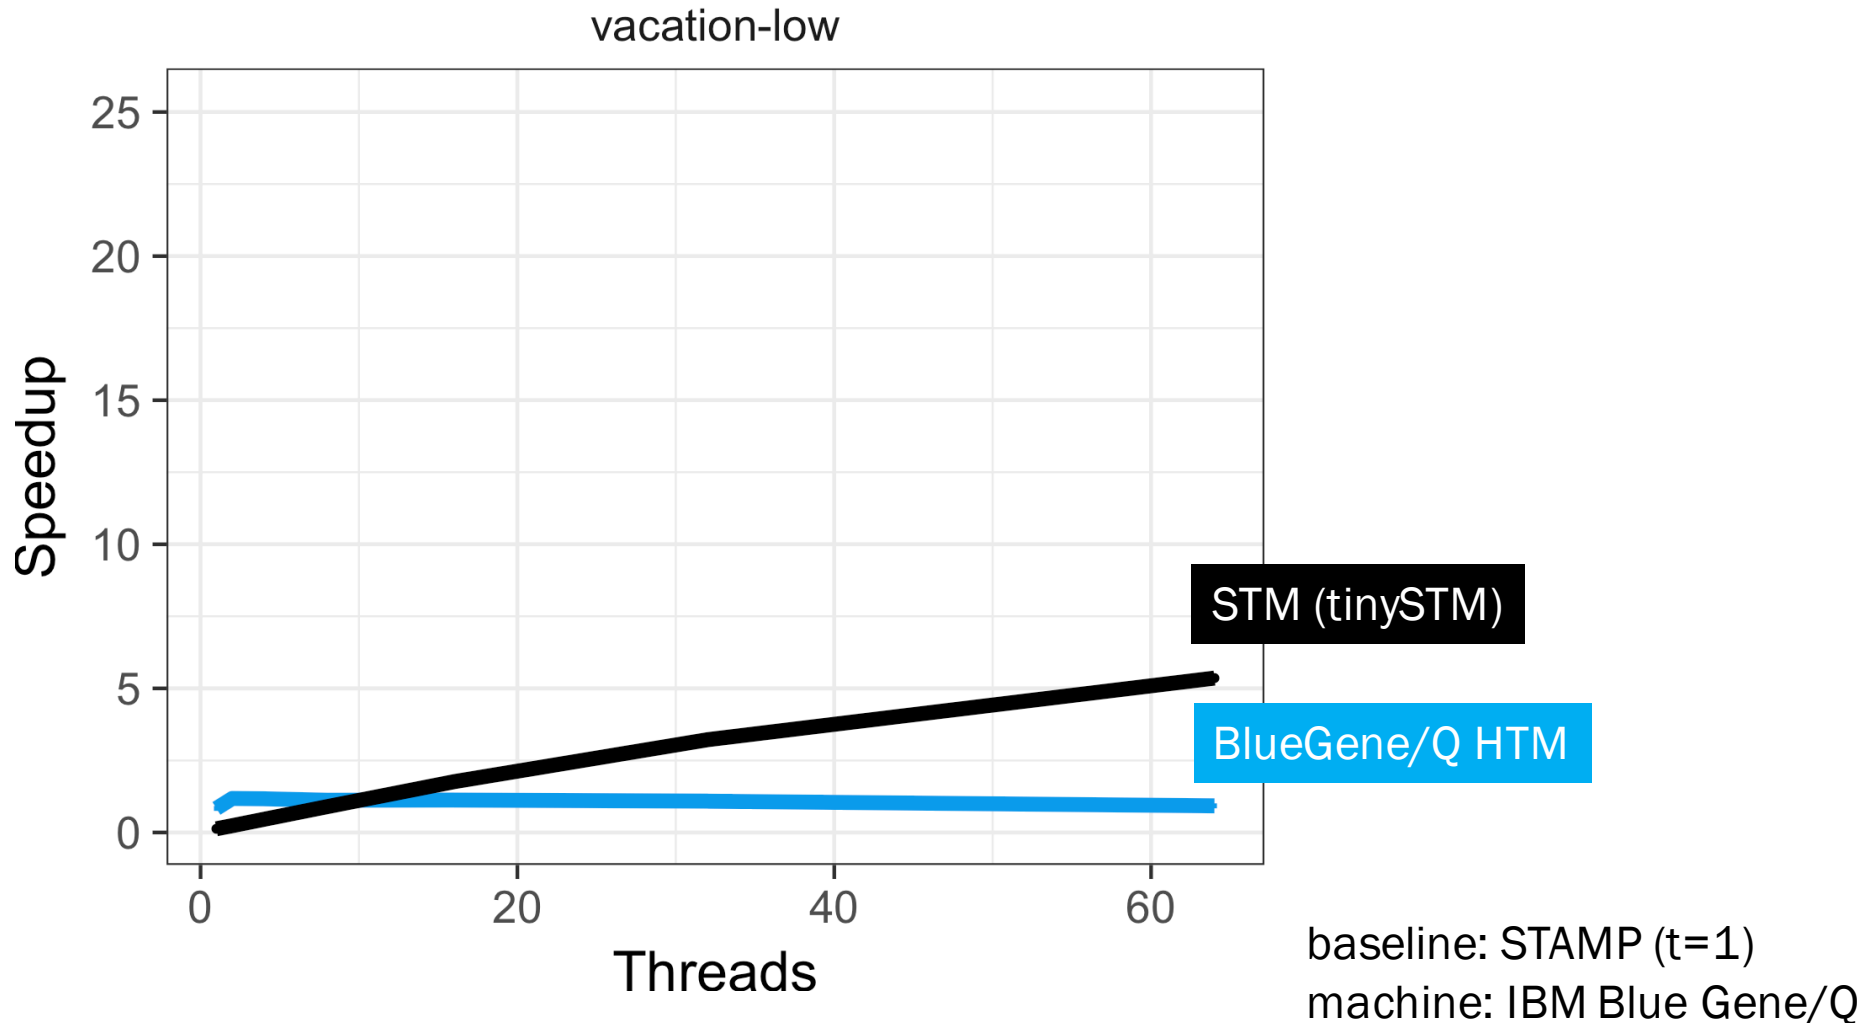
Keshav Pingali

pingali@cs.utexas.edu


Transactional Memory

- Atomic and isolated execution for arbitrary code by tracking memory locations
 - In software (STM) and now in hardware (HTM)
- Studies typically use microbenchmarks or STAMP suite
- Topic of active research for a decade

The Performance to Date



The Performance to Date



Why?

Conceptually, most STAMP benchmarks should be moderately scalable

- String matching
- Clustering
- Packet filtering
- Graph construction
- Accessing a relational database
- Triangular mesh refinement

Failure lies in STAMP programming model

STAMP Programming Model

```
thread_begin
```

```
...
```

```
txn_begin
```

```
... memory accesses ...
```

```
txn_end
```

```
...
```

```
thread_end
```

STAMP Programming Model


```
thread_begin
...
txn_begin
... memory accesses ...
txn_end
...
thread_end
```

- TM protects access to data structures
 - Application data: data structure may not be scalable
 - Worklists: semantics too strict
- Transactions are tied to threads

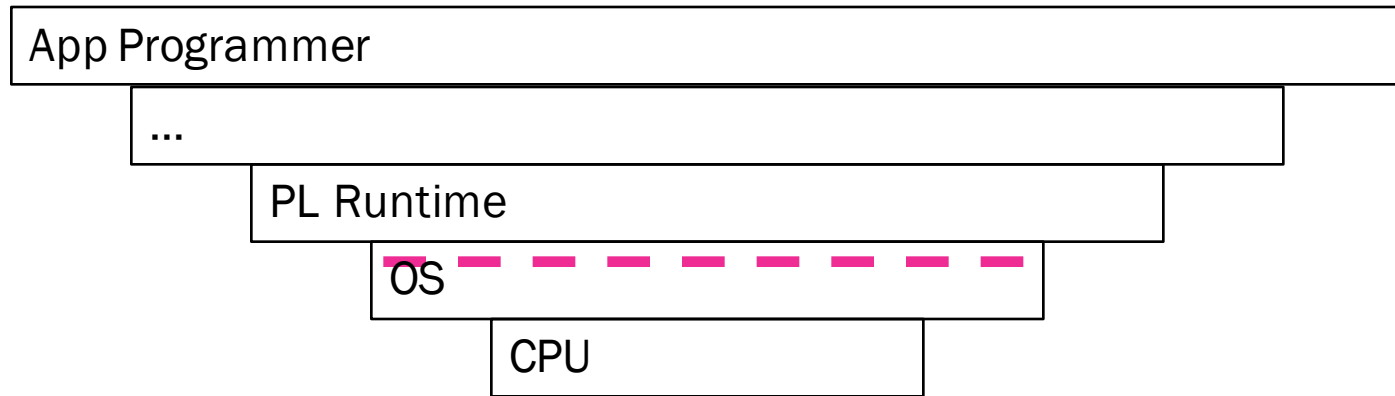
Scalable Programs and Structured Parallelism

App Programmer
...
PL Runtime
OS
CPU

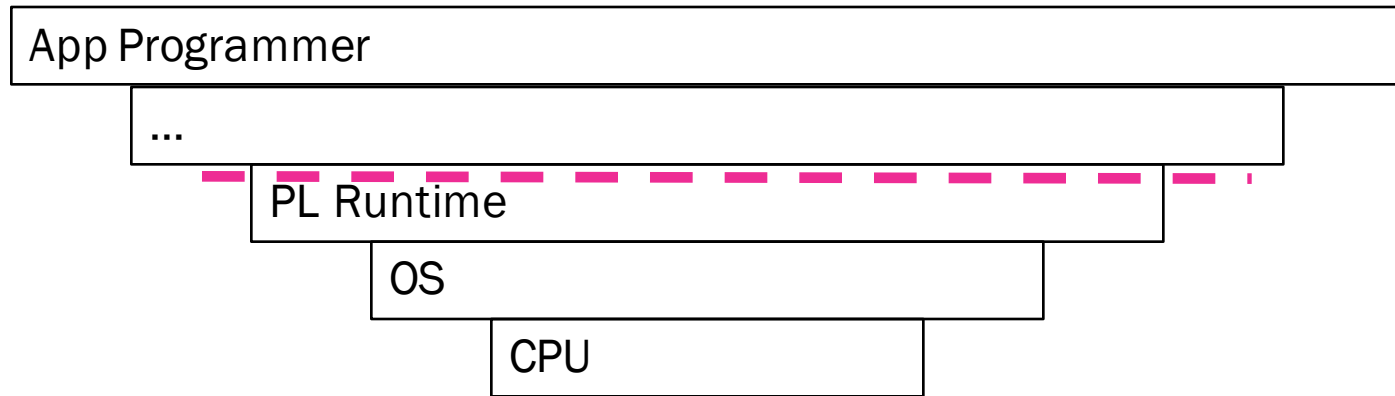
Scalable Programs and Structured Parallelism



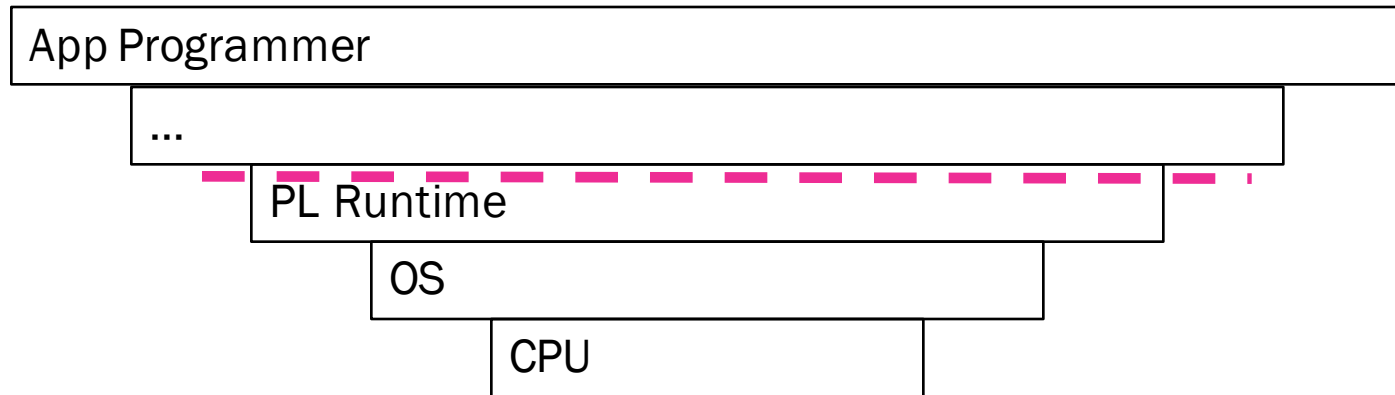
Scalable Programs and Structured Parallelism



Scalable Programs and Structured Parallelism



Scalable Programs and Structured Parallelism



STAMP PM

- Ad-hoc //ism
- Arbitrary data structures

Stampede PM

- Amorphous data //ism [PLDI 2011]
 - Dynamic task generation
 - Tasks may conflict
- Composition of scalable data structures

Stampede Principles

Disjoint Access

Transactions disjoint at logical level should be disjoint at physical level.

Virtualized Transactions

Transactions should be decoupled from threads.

Disjoint Access

Replacing data structures with scalable alternatives

red-black tree => hashtable

hashtable => reduction over hashtables

dynamic buffer => per-task buffer

Disjoint Access

Replacing data structures with scalable alternatives

red-black tree	=>	hashtable
hashtable	=>	reduction over hashtables
dynamic buffer	=>	per-task buffer

Using scheduling-specific data structures

linked-list	=>	workstealing scheduler
counter	=>	workstealing scheduler
dynamic array	=>	workstealing scheduler

Virtualized Transactions

Schedule transactions as tasks

- Execute on any thread in any order
- Exploit runtime information to increase throughput
 - E.g., serialize conflicting tasks
- Use high-performance scheduling data structures

Code changes can be automated

- E.g., compiler support for anonymous functions

XTM

If most transactions access disjoint data, is exclusive conflict detection sufficient?

XTM Implementation

- CAS owner field on transactional access
- Add to list of owned objects
- Reset field on transaction commit or abort

Stampede

Stampede: STAMP transformed to satisfy scalability principles

- Code 88% same as STAMP

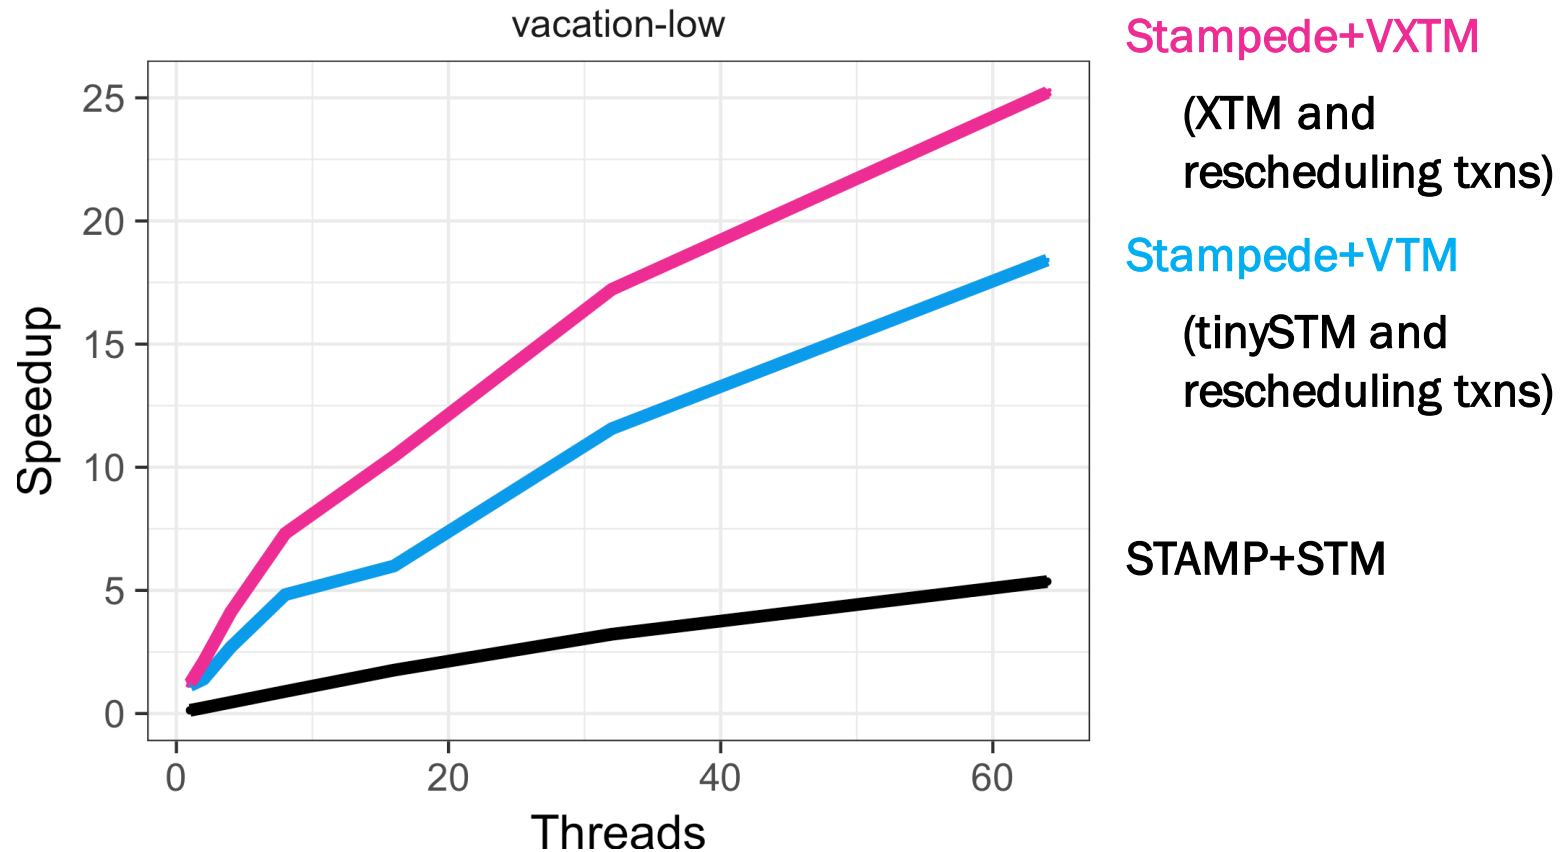
Stampede

Stampede: STAMP transformed to satisfy scalability principles

- Code 88% same as STAMP

Variant	Code	Conflict Detection	Reschedule Transaction?
STAMP+STM	STAMP	tinySTM	No
Stampede+VTM	Stampede	tinySTM	Yes
Stampede+VXTM	Stampede	XTM	Yes

Results




baseline: STAMP (t=1)
machine: Blue Gene/Q

Results

	Median Speedup
STAMP+STM	8.0X
Stampede+VXTM	17.7X

baseline: STAMP (t=1)
machine: Blue Gene/Q

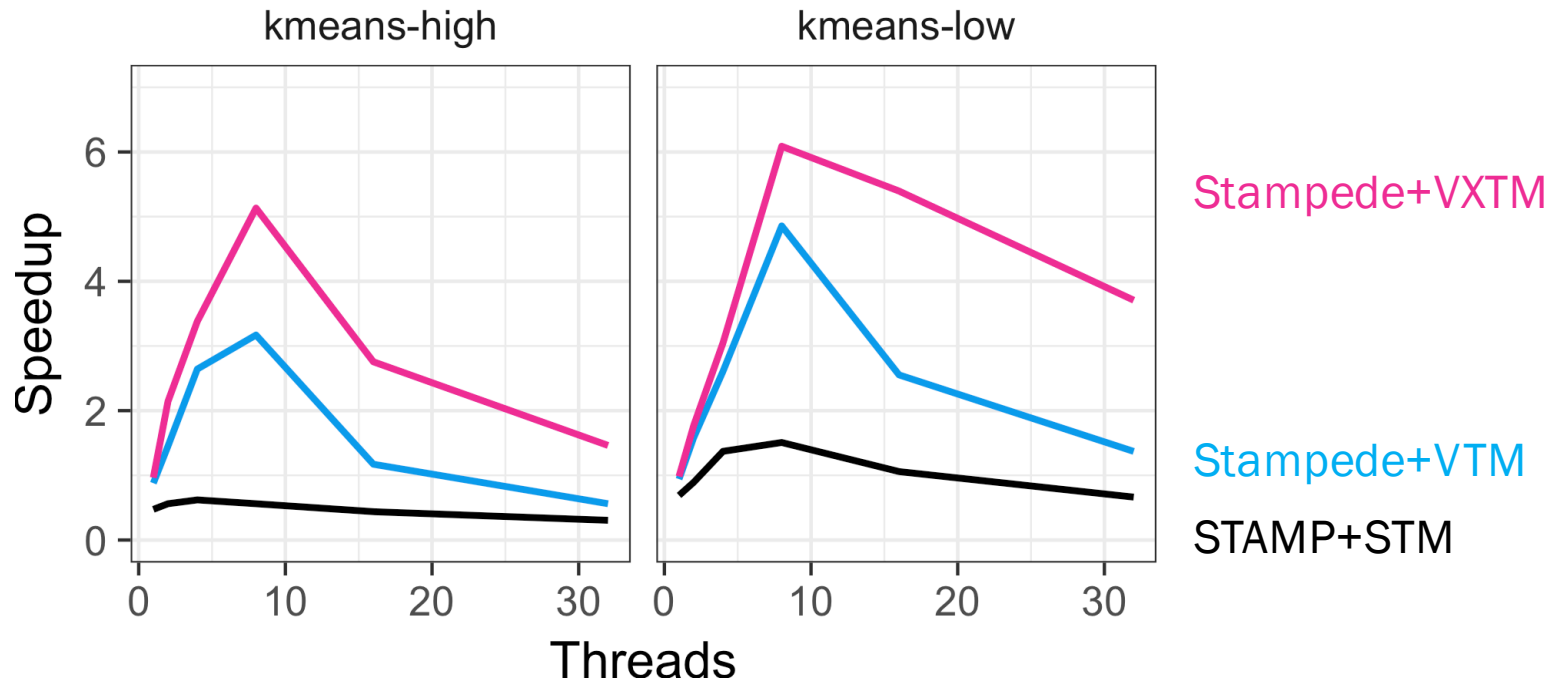
Contention



baseline: STAMP ($t=1$)

machine: Westmere


Contention




- Under high contention, performance still suffers

baseline: STAMP ($t=1$)
machine: Westmere

XTM



XTM



XTM alone is not enough

Conclusion

Stampede+VXTM gives 4X improvement over STAMP+STM/HTM

- Use scalable data structures
 - Reduce conflicts wherever possible
- Exploit scheduling optimizations
 - Do useful work even when conflicts occur
- Simplify TM to VXTM

Is it time to rethink TM interface?

- Division of labor between application, runtime and hardware
- RISC versus CISC

<http://iss.ices.utexas.edu/?p=projects/stampede>